



DML SQL Statements

CE384: Database Design
Maryam Ramezani
Sharif University of Technology
maryam.ramezani@sharif.edu



Introduction

Projection

Table 1

Selection

Table 1

Table 1



Table 2

Introduction

```
CREATE TABLE AGENTS (
    AGENT_CODE CHAR(6) PRIMARY KEY,
    AGENT_NAME CHAR(40),
    WORKING_AREA CHAR(35),
    COMMISSION NUMERIC(10, 2),
    PHONE_NO CHAR(15),
    COUNTRY VARCHAR(25)
);
CREATE TABLE ORDERS (
    ORD_NUM SERIAL PRIMARY KEY,
    ORD_AMOUNT NUMERIC(12, 2) NOT NULL,
    ADVANCE_AMOUNT NUMERIC(12, 2) NOT NULL,
    ORD_DATE DATE NOT NULL,
    CUST_CODE VARCHAR(6) NOT NULL REFERENCES CUSTOMER,
    AGENT_CODE CHAR(6) NOT NULL REFERENCES AGENTS,
    ORD_DESCRIPTION VARCHAR(60) NOT NULL
);
```

```
CREATE TABLE CUSTOMER (
    CUST_CODE VARCHAR(6) PRIMARY KEY,
    CUST_NAME VARCHAR(40) NOT NULL,
    CUST_CITY CHAR(35),
    WORKING_AREA VARCHAR(35) NOT NULL,
    CUST_COUNTRY VARCHAR(20) NOT NULL,
    GRADE INTEGER,
    OPENING_AMT NUMERIC(12, 2) NOT NULL,
    RECEIVE_AMT NUMERIC(12, 2) NOT NULL,
    PAYMENT_AMT NUMERIC(12, 2) NOT NULL,
    OUTSTANDING_AMT NUMERIC(12, 2) NOT NULL,
    PHONE_NO VARCHAR(17) NOT NULL,
    AGENT_CODE CHAR(6) NOT NULL REFERENCES AGENTS
);
```

SELECT

Basic SELECT Statement

```
SELECT      * | { [DISTINCT] column|expression [alias],... }  
FROM        table;
```

- SELECT identifies *what* columns
- FROM identifies *which* table

Selecting All Columns

```
SELECT *  
FROM departments;
```

| DEPARTMENT_ID | DEPARTMENT_NAME | MANAGER_ID | LOCATION_ID |
|---------------|-----------------|------------|-------------|
| 10 | Administration | 200 | 1700 |
| 20 | Marketing | 201 | 1800 |
| 50 | Shipping | 124 | 1500 |
| 60 | IT | 103 | 1400 |
| 80 | Sales | 149 | 2500 |
| 90 | Executive | 100 | 1700 |
| 110 | Accounting | 205 | 1700 |
| 190 | Contracting | | 1700 |

8 rows selected.

Selecting Specific Columns

```
SELECT department_id, location_id  
FROM departments;
```

| DEPARTMENT_ID | LOCATION_ID |
|---------------|-------------|
| 10 | 1700 |
| 20 | 1800 |
| 50 | 1500 |
| 60 | 1400 |
| 80 | 2500 |
| 90 | 1700 |
| 110 | 1700 |
| 190 | 1700 |

8 rows selected.

Arithmetic Expressions

Create expressions with number and date data by using arithmetic operators.

| Operator | Description |
|----------|-------------|
| + | Add |
| - | Subtract |
| * | Multiply |
| / | Divide |

Using Arithmetic Operators

```
SELECT last_name, salary, salary + 300  
FROM employees;
```

| LAST_NAME | SALARY | SALARY+300 |
|-----------|--------|------------|
| King | 24000 | 24300 |
| Kochhar | 17000 | 17300 |
| De Haan | 17000 | 17300 |
| Hunold | 9000 | 9300 |
| Ernst | 6000 | 6300 |

...

| | | |
|-----------|-------|-------|
| Hartstein | 13000 | 13300 |
| Fay | 6000 | 6300 |
| Higgins | 12000 | 12300 |
| Gietz | 8300 | 8600 |

20 rows selected.

Operator Precedence & Using Parentheses

```
SELECT last_name, salary, 12*salary+100  
FROM employees;
```

| LAST_NAME | SALARY | 12*SALARY+100 |
|-----------|--------|---------------|
| King | 24000 | 288100 |
| Kochhar | 17000 | 204100 |
| De Haan | 17000 | 204100 |
| Hunold | 9000 | 108100 |
| Ernst | 6000 | 72100 |
| ... | | |
| Hartstein | 13000 | 156100 |
| Fay | 6000 | 72100 |
| Higgins | 12000 | 144100 |
| Gietz | 8300 | 99700 |

20 rows selected.

```
SELECT last_name, salary, 12*(salary+100)  
FROM employees;
```

| LAST_NAME | SALARY | 12*(SALARY+100) |
|-----------|--------|-----------------|
| King | 24000 | 289200 |
| Kochhar | 17000 | 205200 |
| De Haan | 17000 | 205200 |
| Hunold | 9000 | 109200 |
| Ernst | 6000 | 73200 |
| ... | | |
| Hartstein | 13000 | 157200 |
| Fay | 6000 | 73200 |
| Higgins | 12000 | 145200 |
| Gietz | 8300 | 100800 |

20 rows selected.

Defining a Null Value

- A null is a value that is unavailable, unassigned, unknown, or inapplicable.
- A null is not the same as zero or a blank space.
- Arithmetic expressions containing a null value evaluate to null.

```
SELECT last_name, 12*salary*commission_pct  
FROM employees;
```

| LAST_NAME | JOB_ID | SALARY | COMMISSION_PCT |
|-----------|---------|--------|----------------|
| King | AD_PRES | 24000 | |
| Kochhar | AD_VP | 17000 | |

| | | | |
|---------|--------|-------|----|
| Zlotkey | SA_MAN | 10500 | .2 |
| Abel | SA_REP | 11000 | .3 |
| Taylor | SA_REP | 8600 | .2 |

| | | | |
|-------|------------|------|--|
| Gietz | AC_ACCOUNT | 8300 | |
|-------|------------|------|--|

20 rows selected.

Using Column Aliases

A column alias:

- Renames a column heading
- Is useful with calculations
- Immediately follows the column name - there can also be the optional AS keyword between the column name and alias
- Requires double quotation marks if it contains spaces or special characters or is

case sensitive

```
SELECT last_name AS name, commission_pct comm  
FROM employees;
```

| NAME | COMM |
|---------|------|
| King | |
| Kochhar | |
| De Haan | |
| ... | |

20 rows selected.

```
SELECT last_name "Name", salary*12 "Annual Salary"  
FROM employees;
```

| Name | Annual Salary |
|---------|---------------|
| King | 288000 |
| Kochhar | 204000 |
| De Haan | 204000 |
| ... | |

20 rows selected.

Concatenation Operator

A concatenation operator:

- Concatenates columns or character strings to other columns
- Is represented by two vertical bars (||)
- Creates a resultant column that is a character expression

```
SELECT      last_name || job_id AS "Employees"  
FROM        employees;
```

| Employees |
|----------------|
| KingAD_PRES |
| KochharAD_VP |
| De HaanAD_VP |
| HunoldIT_PROG |
| ErnstIT_PROG |
| LorentzIT_PROG |
| MourgosST_MAN |
| RajsST_CLERK |
| ... |

20 rows selected.

Literal Character Strings

- A literal is a character, a number, or a date included in the SELECT list.
- Date and character literal values must be enclosed within single quotation marks.
- Each character string is output once for each row returned.

```
SELECT last_name      ||' is a '|| job_id  
      AS "Employee Details"  
  FROM employees;
```

| Employee Details |
|----------------------|
| King is a AD_PRES |
| Kochhar is a AD_VP |
| De Haan is a AD_VP |
| Hunold is a IT_PROG |
| Ernst is a IT_PROG |
| Lorentz is a IT_PROG |
| Mourgos is a ST_MAN |
| Rajs is a ST_CLERK |
| ... |

20 rows selected.

Duplicate Rows

- The default display of queries is all rows, including duplicate rows.

```
SELECT department_id  
FROM employees;
```

| DEPARTMENT_ID | DEPARTMENT_NAME |
|---------------|--------------------|
| 90 | Customer Service |
| 90 | Marketing |
| 90 | Human Resources |
| 60 | Sales |
| 60 | Product Management |
| 60 | R&D |
| 50 | Finance |
| 50 | Operations |
| 50 | Infrastructure |

20 rows selected.

Eliminating Duplicate Rows

- Eliminate duplicate rows by using the DISTINCT keyword in the SELECT clause.

```
SELECT DISTINCT department_id  
FROM employees;
```

| DEPARTMENT_ID |
|---------------|
| 10 |
| 20 |
| 50 |
| 60 |
| 80 |
| 90 |
| 110 |
| |

8 rows selected.

Restricting and Sorting

Limiting Rows Using a Selection

EMPLOYEES

| EMPLOYEE_ID | LAST_NAME | JOB_ID | DEPARTMENT_ID |
|-------------|-----------|---------|---------------|
| 100 | King | AD_PRES | 90 |
| 101 | Kochhar | AD_VP | 90 |
| 102 | De Haan | AD_VP | 90 |
| 103 | Hunold | IT_PROG | 60 |
| 104 | Ernst | IT_PROG | 60 |
| 107 | Lorentz | IT_PROG | 60 |
| 124 | Mourgos | ST_MAN | 50 |
| ... | | | |

20 rows selected.

**“retrieve all
employees
in department 90”**



| EMPLOYEE_ID | LAST_NAME | JOB_ID | DEPARTMENT_ID |
|-------------|-----------|---------|---------------|
| 100 | King | AD_PRES | 90 |
| 101 | Kochhar | AD_VP | 90 |
| 102 | De Haan | AD_VP | 90 |

Limiting the Rows Selected

- Restrict the rows returned by using the WHERE clause.
- The WHERE clause follows the FROM clause.

```
SELECT      * | { [DISTINCT] column|expression [alias],... }  
FROM        table  
[WHERE      condition(s)];
```

Using the WHERE Clause

- With fixed value

```
SELECT employee_id, last_name, job_id, department_id  
FROM   employees  
WHERE  department_id = 90 ;
```

| EMPLOYEE_ID | LAST_NAME | JOB_ID | DEPARTMENT_ID |
|-------------|-----------|---------|---------------|
| 100 | King | AD_PRES | 90 |
| 101 | Kochhar | AD_VP | 90 |
| 102 | De Haan | AD_VP | 90 |

- With variable

```
SELECT employee_id, last_name, job_id, department_id  
FROM   employees  
WHERE  department_id = :input ;
```

Character Strings and Dates

- Character strings and date values are enclosed in single quotation marks.
- Character values are case sensitive, and date values are format sensitive.
- The default date format is DD-MON-RR.

```
SELECT last_name, job_id, department_id  
FROM   employees  
WHERE  last_name = 'Whalen';
```

Comparison Conditions

| Operator | Meaning |
|----------|--------------------------|
| = | Equal to |
| > | Greater than |
| >= | Greater than or equal to |
| < | Less than |
| <= | Less than or equal to |
| ▷ | Not equal to |

| Operator | Meaning |
|----------------------|---------------------------------|
| BETWEEN ...AND... | Between two values (inclusive), |
| IN(set) | Match any of a list of values |
| LIKE | Match a character pattern |
| IS NULL | Is a null value |

Using Comparison Conditions

```
SELECT last_name, salary  
FROM   employees  
WHERE  salary <= 3000;
```

| LAST_NAME | SALARY |
|-----------|--------|
| Matos | 2600 |
| Vargas | 2500 |

Using the BETWEEN Condition

- Use the BETWEEN condition to display rows based on a range of values.

```
SELECT last_name, salary  
FROM employees  
WHERE salary BETWEEN 2500 AND 3500;
```

Lower limit Upper limit

| LAST_NAME | SALARY |
|-----------|--------|
| Rajs | 3500 |
| Davies | 3100 |
| Matos | 2600 |
| Vargas | 2500 |

Using the IN Condition

- Use the IN membership condition to test for values in a list.

```
SELECT employee_id, last_name, salary, manager_id  
FROM   employees  
WHERE  manager_id IN (100, 101, 201);
```

| EMPLOYEE_ID | LAST_NAME | SALARY | MANAGER_ID |
|-------------|-----------|--------|------------|
| 202 | Fay | 6000 | 201 |
| 200 | Whalen | 4400 | 101 |
| 205 | Higgins | 12000 | 101 |
| 101 | Kochhar | 17000 | 100 |
| 102 | De Haan | 17000 | 100 |
| 124 | Mourgos | 5800 | 100 |
| 149 | Zlotkey | 10500 | 100 |
| 201 | Hartstein | 13000 | 100 |

8 rows selected.

Using the LIKE Condition

- Use the LIKE condition to perform wildcard searches of valid search string values.
- Search conditions can contain either literal characters or numbers:
 - % denotes zero or many characters.
 - _ denotes one character.

```
SELECT      first_name
FROM        employees
WHERE       first_name LIKE 'S%';
```

Using the LIKE Condition

- You can combine pattern-matching characters.

```
SELECT last_name  
FROM   employees  
WHERE  last_name LIKE '_o%';
```

| LAST_NAME |
|-----------|
| Kochhar |
| Lorentz |
| Mourgos |

- You can use the ESCAPE identifier to search for the actual % and _ symbols.

Using the NULL Conditions

- Test for nulls with the IS NULL operator.

```
SELECT last_name, manager_id  
FROM employees  
WHERE manager_id IS NULL;
```

| LAST_NAME | MANAGER_ID |
|-----------|------------|
| King | |

Logical Conditions

| Operator | Meaning |
|----------|---|
| AND | Returns TRUE if <i>both</i> component conditions are true |
| OR | Returns TRUE if <i>either</i> component condition is true |
| NOT | Returns TRUE if the following condition is false |

Using the AND Operator

- AND requires both conditions to be true.

```
SELECT employee_id, last_name, job_id, salary  
FROM employees  
WHERE salary >=10000  
AND job_id LIKE '%MAN%';
```

| EMPLOYEE_ID | LAST_NAME | JOB_ID | SALARY |
|-------------|-----------|--------|--------|
| 149 | Zlotkey | SA_MAN | 10500 |
| 201 | Hartstein | MK_MAN | 13000 |

Using the OR Operator

- OR requires either condition to be true.

```
SELECT employee_id, last_name, job_id, salary  
FROM employees  
WHERE salary >= 10000  
OR job_id LIKE '%MAN%';
```

| EMPLOYEE_ID | LAST_NAME | JOB_ID | SALARY |
|-------------|-----------|---------|--------|
| 100 | King | AD_PRES | 24000 |
| 101 | Kochhar | AD_VP | 17000 |
| 102 | De Haan | AD_VP | 17000 |
| 124 | Mourgos | ST_MAN | 5800 |
| 149 | Zlotkey | SA_MAN | 10500 |
| 174 | Abel | SA_REP | 11000 |
| 201 | Hartstein | MK_MAN | 13000 |
| 205 | Higgins | AC_MGR | 12000 |

8 rows selected.

Using the NOT Operator

```
SELECT last_name, job_id  
FROM employees  
WHERE job_id  
NOT IN ('IT_PROG', 'ST_CLERK', 'SA_REP');
```

| LAST_NAME | JOB_ID |
|-----------|------------|
| King | AD_PRES |
| Kochhar | AD_VP |
| De Haan | AD_VP |
| Mourgos | ST_MAN |
| Zlotkey | SA_MAN |
| Whalen | AD_ASST |
| Hartstein | MK_MAN |
| Fay | MK_REP |
| Higgins | AC_MGR |
| Gietz | AC_ACCOUNT |

10 rows selected.

Rules of Precedence

| Order Evaluated | Operator |
|-----------------|-------------------------------|
| 1 | Arithmetic operators |
| 2 | Concatenation operator |
| 3 | Comparison conditions |
| 4 | IS [NOT] NULL, LIKE, [NOT] IN |
| 5 | [NOT] BETWEEN |
| 6 | NOT logical condition |
| 7 | AND logical condition |
| 8 | OR logical condition |

Override rules of precedence by using parentheses.

Rules of Precedence

```
SELECT last_name, job_id, salary  
FROM   employees  
WHERE  job_id = 'SA_REP'  
OR     job_id = 'AD_PRES'  
AND    salary > 15000;
```

| LAST_NAME | JOB_ID | SALARY |
|-----------|---------|--------|
| King | AD_PRES | 24000 |
| Abel | SA_REP | 11000 |
| Taylor | SA_REP | 8600 |
| Grant | SA_REP | 7000 |

Rules of Precedence

- Use parentheses to force priority.

```
SELECT last_name, job_id, salary
FROM   employees
WHERE  (job_id = 'SA_REP'
OR      job_id = 'AD_PRES')
AND    salary > 15000;
```

| LAST_NAME | JOB_ID | SALARY |
|-----------|---------|--------|
| King | AD_PRES | 24000 |

ORDER BY Clause

- Sort rows with the ORDER BY clause
 - ASC: ascending order, default
 - DESC: descending order
- The ORDER BY clause comes last in the SELECT statement.

```
SELECT last_name, job_id, department_id, hire_date  
FROM employees  
ORDER BY hire_date ;
```

| LAST_NAME | JOB_ID | DEPARTMENT_ID | HIRE_DATE |
|-----------|---------|---------------|-----------|
| King | AD_PRES | 90 | 17-JUN-87 |
| Whalen | AD_ASST | 10 | 17-SEP-87 |
| Kochhar | AD_VP | 90 | 21-SEP-89 |
| Hunold | IT_PROG | 60 | 03-JAN-90 |
| Ernst | IT_PROG | 60 | 21-MAY-91 |

...

20 rows selected.

Sorting in Descending Order

```
SELECT      last_name, job_id, department_id, hire_date  
FROM        employees  
ORDER BY    hire_date DESC ;
```

| LAST_NAME | JOB_ID | DEPARTMENT_ID | HIRE_DATE |
|-----------|----------|---------------|-----------|
| Zlotkey | SA_MAN | 80 | 29-JAN-00 |
| Mourgos | ST_MAN | 50 | 16-NOV-99 |
| Grant | SA_REP | | 24-MAY-99 |
| Lorentz | IT_PROG | 60 | 07-FEB-99 |
| Vargas | ST_CLERK | 50 | 09-JUL-98 |
| Taylor | SA_REP | 80 | 24-MAR-98 |
| Matos | ST_CLERK | 50 | 15-MAR-98 |
| Fay | MK_REP | 20 | 17-AUG-97 |
| Davies | ST_CLERK | 50 | 29-JAN-97 |
| ... | | | |

20 rows selected.

Sorting by Column Alias

```
SELECT employee_id, last_name, salary*12 annsal  
FROM employees  
ORDER BY annsal;
```

| EMPLOYEE_ID | LAST_NAME | ANNSAL |
|-------------|-----------|--------|
| 144 | Vargas | 30000 |
| 143 | Matos | 31200 |
| 142 | Davies | 37200 |
| 141 | Rajs | 42000 |
| 107 | Lorentz | 50400 |
| 200 | Whalen | 52800 |
| 124 | Mourgos | 69600 |
| 104 | Ernst | 72000 |
| 202 | Fay | 72000 |
| 178 | Grant | 84000 |
| ... | | |

20 rows selected.

Sorting by Multiple Columns

- The order of ORDER BY list is the order of sort.

```
SELECT last_name, department_id, salary  
FROM employees  
ORDER BY department_id, salary DESC;
```

| LAST_NAME | DEPARTMENT_ID | SALARY |
|-----------|---------------|--------|
| Whalen | 10 | 4400 |
| Hartstein | 20 | 13000 |
| Fay | 20 | 6000 |
| Mourgos | 50 | 5800 |
| Rajs | 50 | 3500 |
| Davies | 50 | 3100 |
| Matos | 50 | 2600 |
| Vargas | 50 | 2500 |
| ... | | |

20 rows selected.

- You can sort by a column that is not in the SELECT list.

Summary

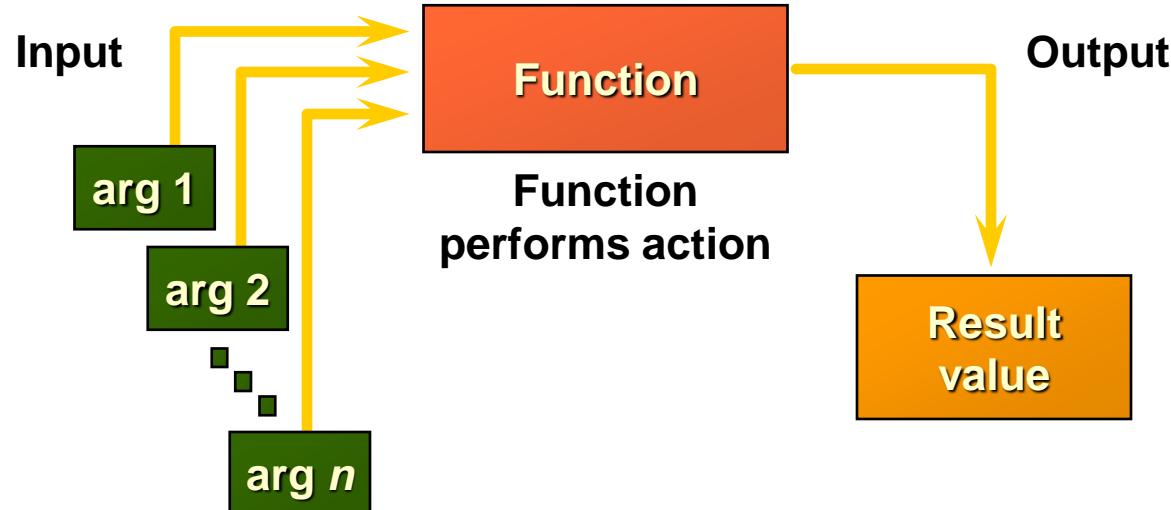
In this lesson, you should have learned how to:

- Use the WHERE clause to restrict rows of output
 - Use the comparison conditions
 - Use the BETWEEN, IN, LIKE, and NULL conditions
 - Apply the logical AND, OR, and NOT operators
- Use the ORDER BY clause to sort rows of output

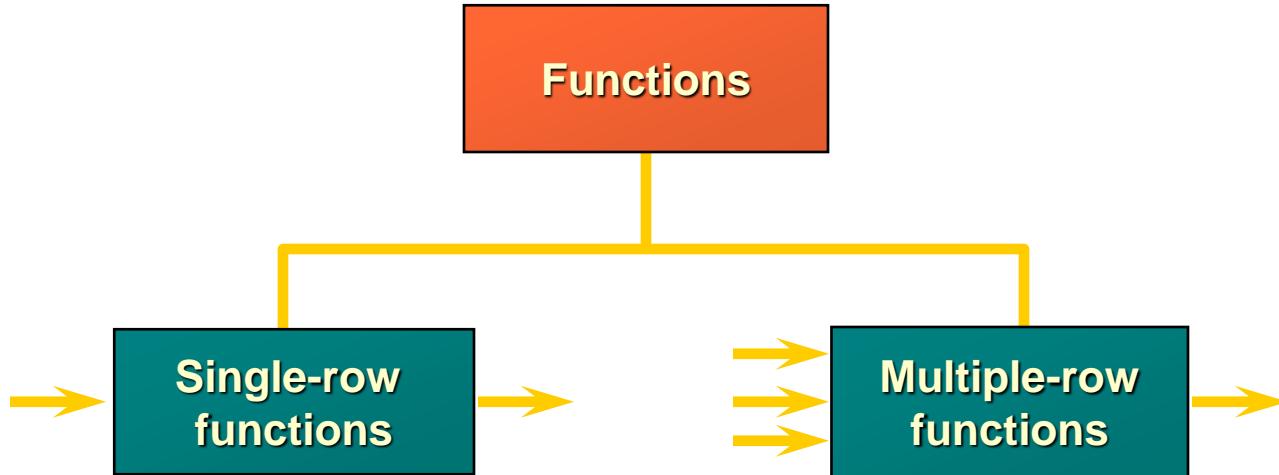
```
SELECT      * | { [DISTINCT] column|expression [alias], ... }  
FROM        table  
[WHERE      condition(s)]  
[ORDER BY   {column, expr, alias} [ASC|DESC]] ;
```

Single-Row Functions

SQL Functions



Two Types of SQL Functions



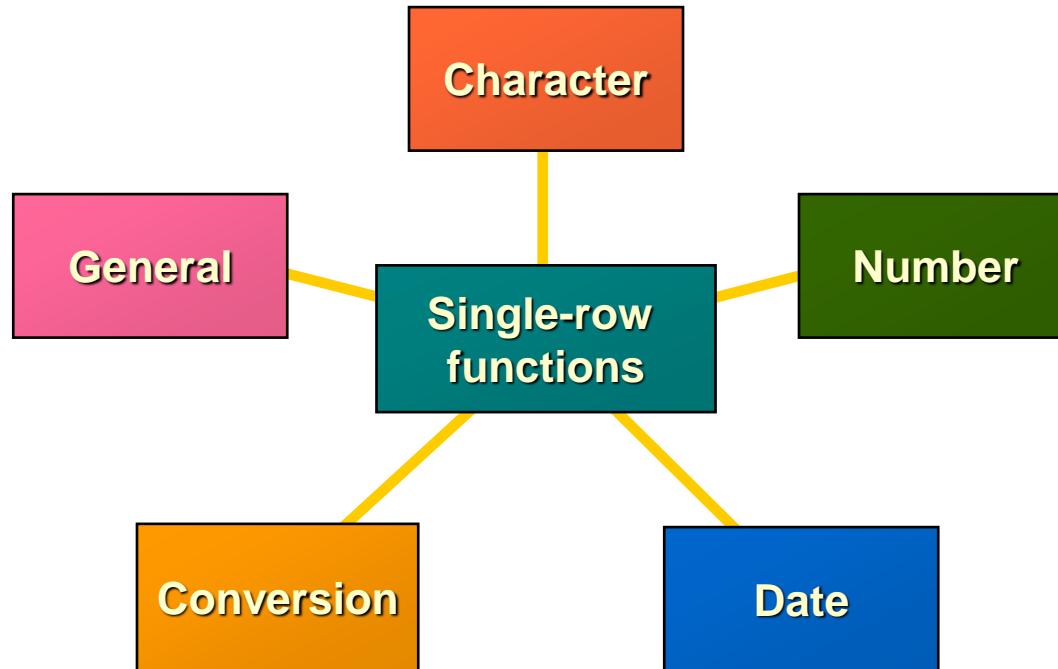
Single-Row Functions

Single row functions:

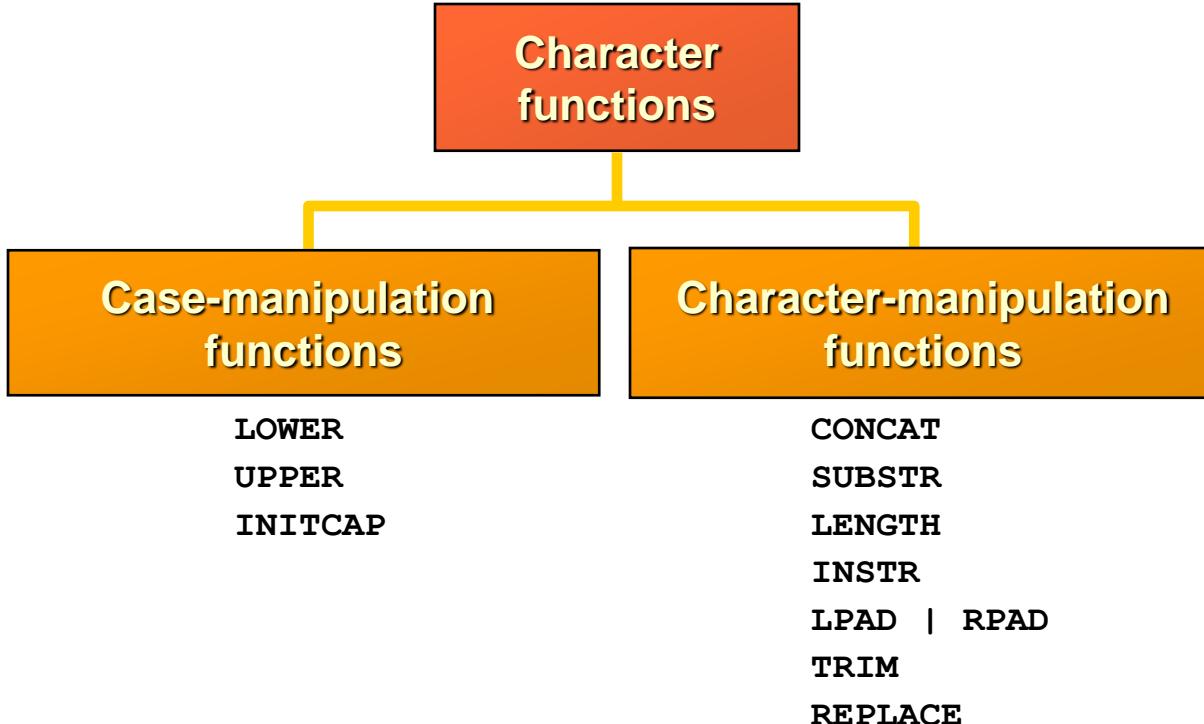
- Manipulate data items
- Accept arguments and return one value
- Act on each row returned
- Return one result per row
- May modify the data type
- Can be nested
- Accept arguments which can be a column or an expression

function_name [(arg1, arg2, ...)]

Single-Row Functions



Character Functions



Case Manipulation Functions

- These functions convert case for character strings.

| Function | Result |
|-----------------------|------------|
| LOWER('SQL Course') | sql course |
| UPPER('SQL Course') | SQL COURSE |
| INITCAP('SQL Course') | Sql Course |

Using Case Manipulation Functions

- Display the employee number, name, and department number for employee Higgins:

```
SELECT employee_id, last_name, department_id  
FROM   employees  
WHERE  last_name = 'higgins';  
no rows selected
```

```
SELECT employee_id, last_name, department_id  
FROM   employees  
WHERE  LOWER(last_name) = 'higgins';
```

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID |
|-------------|-----------|---------------|
| 205 | Higgins | 110 |

Character-Manipulation Functions

- These functions manipulate character strings:

| Function | Result |
|-----------------------------|------------|
| CONCAT('Hello', 'World') | HelloWorld |
| SUBSTR('HelloWorld',1,5) | Hello |
| LENGTH('HelloWorld') | 10 |
| INSTR('HelloWorld', 'W') | 6 |
| LPAD(salary,10,'*') | *****24000 |
| RPAD(salary, 10, '*') | 24000***** |
| TRIM('H' FROM 'HelloWorld') | elloWorld |

Using the Character-Manipulation Functions

```
SELECT employee_id, CONCAT(first_name, last_name) NAME,  
      job_id, LENGTH(last_name),  
      INSTR(last_name, 'a') "Contains 'a'?"  
FROM employees  
WHERE SUBSTR(job_id, 4) = 'REP';
```

1

2

3

| EMPLOYEE_ID | NAME | JOB_ID | LENGTH(LAST_NAME) | Contains 'a?' |
|-------------|----------------|--------|-------------------|---------------|
| 174 | EllenAbel | SA_REP | 4 | 0 |
| 176 | JonathonTaylor | SA_REP | 6 | 2 |
| 178 | KimberelyGrant | SA_REP | 5 | 3 |
| 202 | PatFay | MK_REP | 3 | 2 |

1

2

3

Number Functions

- ROUND: Rounds value to specified decimal

ROUND(45.926, 2)  45.93

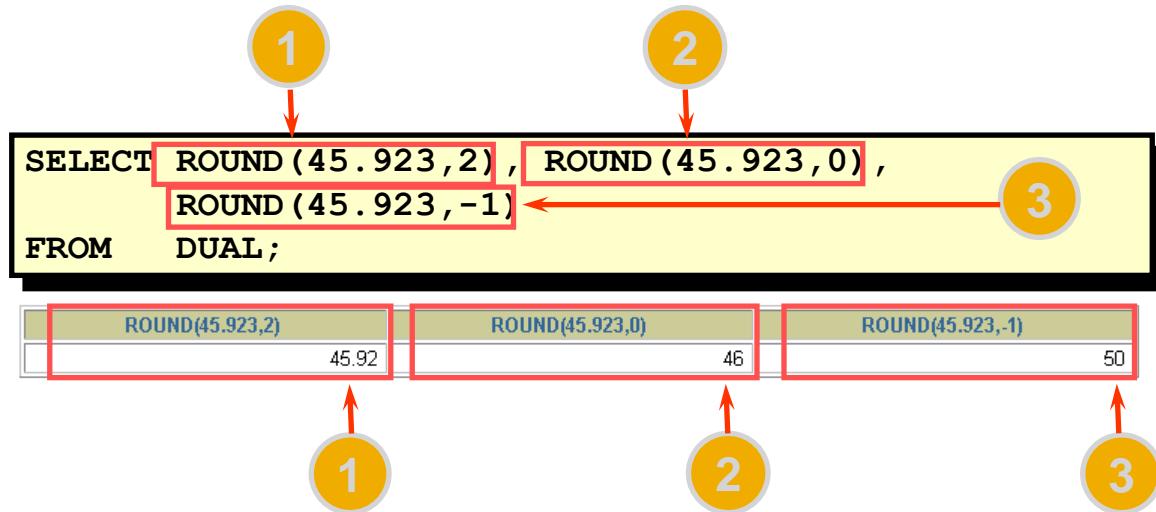
- TRUNC: Truncates value to specified decimal

TRUNC(45.926, 2)  45.92

- MOD: Returns remainder of division

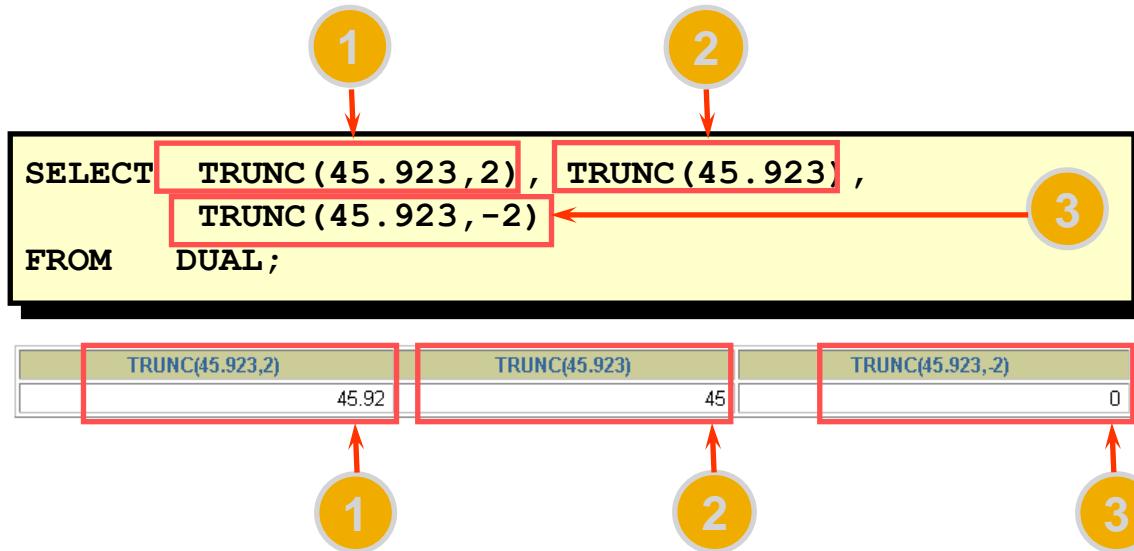
MOD(1600, 300)  100

Using the ROUND Function



DUAL is a dummy table you can use to view results from functions and calculations. Postgres does not need it!!

Using the TRUNC Function



Using the MOD Function

- Calculate the remainder of a salary after it is divided by 5000 for all employees whose job title is sales representative.

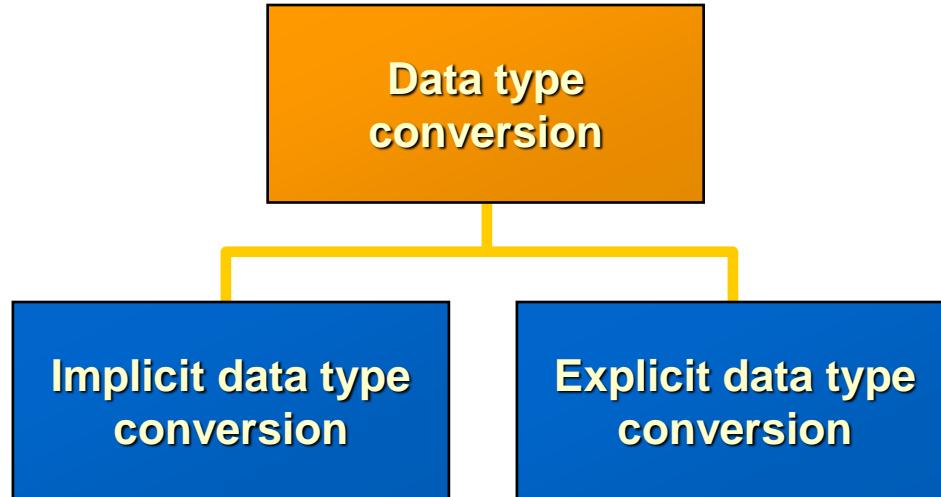
```
SELECT last_name, salary, MOD(salary, 5000)  
FROM   employees  
WHERE  job_id = 'SA_REP';
```

| LAST_NAME | SALARY | MOD(SALARY,5000) |
|-----------|--------|------------------|
| Abel | 11000 | 1000 |
| Taylor | 8600 | 3600 |
| Grant | 7000 | 2000 |

Working with Dates

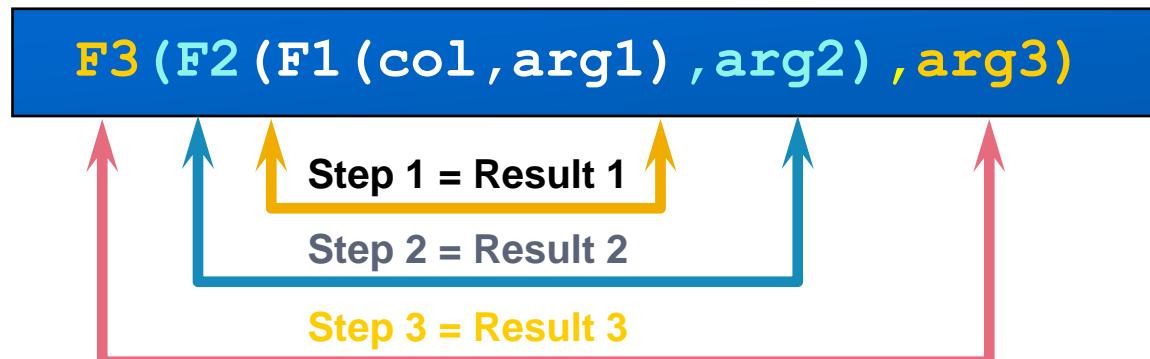
| Function | Description |
|-----------------------------|------------------------------------|
| <code>MONTHS_BETWEEN</code> | Number of months between two dates |
| <code>ADD_MONTHS</code> | Add calendar months to date |
| <code>NEXT_DAY</code> | Next day of the date specified |
| <code>LAST_DAY</code> | Last day of the month |
| <code>ROUND</code> | Round date |
| <code>TRUNC</code> | Truncate date |

Conversion Functions



Nesting Functions

- Single-row functions can be nested to any level.
- Nested functions are evaluated from deepest level to the least deep level.



Nesting Functions

```
SELECT last_name,  
       coalesce (null,null,'No Manager')  
FROM   employees  
WHERE  manager_id IS NULL;
```

| LAST_NAME | NVL(TO_CHAR(MANAGER_ID),'NOMANAGER') |
|-----------|--------------------------------------|
| King | No Manager |

Using the CASE Expression

- Facilitates conditional inquiries by doing the work of an IF-THEN-ELSE statement:

```
SELECT last_name, job_id, salary,  
       CASE job_id WHEN 'IT_PROG' THEN 1.10*salary  
                     WHEN 'ST_CLERK' THEN 1.15*salary  
                     WHEN 'SA REP' THEN 1.20*salary  
                     ELSE salary END      "REVISED_SALARY"  
FROM employees;
```

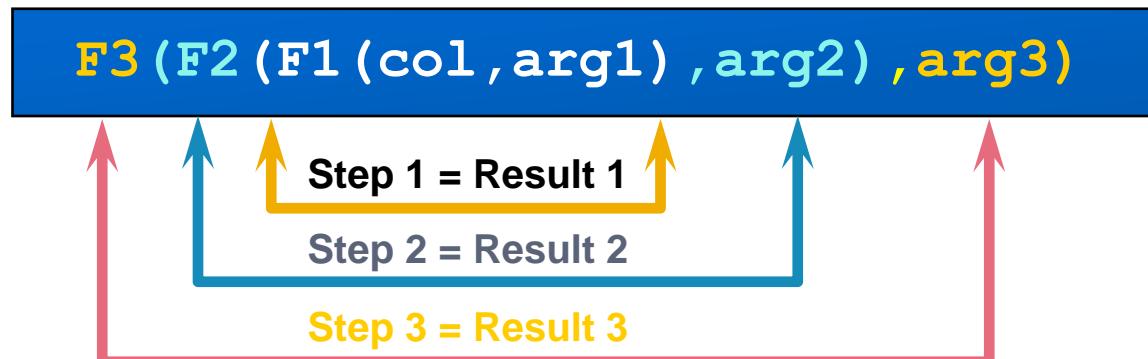
| LAST_NAME | JOB_ID | SALARY | REVISED_SALARY |
|-----------|------------|--------|----------------|
| Lorentz | IT_PROG | 4200 | 4620 |
| Mourgos | ST_MAN | 5800 | 5800 |
| Rajs | ST_CLERK | 3500 | 4025 |
| Gietz | AC_ACCOUNT | 8300 | 8300 |

20 rows selected.

Displaying Data from Multiple Tables

Nesting Functions

- Single-row functions can be nested to any level.
- Nested functions are evaluated from deepest level to the least deep level.



Displaying Data from Multiple Tables

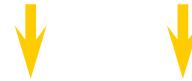
Obtaining Data from Multiple Tables

EMPLOYEES

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID |
|-------------|-----------|---------------|
| 100 | King | 90 |
| 101 | Kochhar | 90 |
| ... | | |
| 202 | Fay | 20 |
| 205 | Higgins | 110 |
| 206 | Gietz | 110 |

DEPARTMENTS

| DEPARTMENT_ID | DEPARTMENT_NAME | LOCATION_ID |
|---------------|-----------------|-------------|
| 10 | Administration | 1700 |
| 20 | Marketing | 1800 |
| 50 | Shipping | 1500 |
| 60 | IT | 1400 |
| 80 | Sales | 2500 |
| 90 | Executive | 1700 |
| 110 | Accounting | 1700 |
| 190 | Contracting | 1700 |



| EMPLOYEE_ID | DEPARTMENT_ID | DEPARTMENT_NAME |
|-------------|---------------|-----------------|
| 200 | 10 | Administration |
| 201 | 20 | Marketing |
| 202 | 20 | Marketing |
| ... | | |
| 102 | 90 | Executive |
| 205 | 110 | Accounting |
| 206 | 110 | Accounting |

Cartesian Products

- A join combines two or more tables side by side. If you do not specify how to join the tables, you get a Cartesian product. This means that SQL combines each row from the first table with every row from the second table.
- A Cartesian product is formed when:
 - A join condition is omitted
 - A join condition is invalid
 - All rows in the first table are joined to all rows in the second table
- To avoid a Cartesian product, always include a valid join condition in a WHERE clause.

Generating a Cartesian Product

- SELECT A.* , B.* FROM FRUITS A, SIZES B

| Fruits |
|---------|
| Apples |
| Mangoes |

| Sizes |
|--------|
| Small |
| Medium |
| Big |

Cartesian Product And Resultant Data

| Fruits | Sizes |
|---------|--------|
| Apples | Small |
| Mangoes | Small |
| Apples | Medium |
| Mangoes | Medium |
| Apples | Big |
| Mangoes | Big |

Generating a Cartesian Product

EMPLOYEES (20 rows)

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID |
|-------------|-----------|---------------|
| 100 | King | 90 |
| 101 | Kochhar | 90 |
| ... | | |
| 202 | Fay | 20 |
| 205 | Higgins | 110 |
| 206 | Gietz | 110 |

20 rows selected.

DEPARTMENTS (8 rows)

| DEPARTMENT_ID | DEPARTMENT_NAME | LOCATION_ID |
|---------------|-----------------|-------------|
| 10 | Administration | 1700 |
| 20 | Marketing | 1800 |
| 50 | Shipping | 1500 |
| 60 | IT | 1400 |
| 80 | Sales | 2500 |
| 90 | Executive | 1700 |
| 110 | Accounting | 1700 |
| 190 | Contracting | 1700 |

8 rows selected.

Cartesian
product:
 $20 \times 8 = 160$ rows



| EMPLOYEE_ID | DEPARTMENT_ID | LOCATION_ID |
|-------------|---------------|-------------|
| 100 | 90 | 1700 |
| 101 | 90 | 1700 |
| 102 | 90 | 1700 |
| 103 | 60 | 1700 |
| 104 | 60 | 1700 |
| 107 | 60 | 1700 |
| ... | | |

160 rows selected.

Join

- SELECT A.fruitName, B.sizeName FROM FRUITS A,SIZES B WHERE A.FRUITID = B.FRUITID;

| | fruitName | sizeName |
|---|-----------|----------|
| 1 | Apples | Small |
| 2 | Apples | Big |
| 3 | Mangoes | Medium |

If we apply the join condition, we will get the output accordingly as given here. In this way, we can avoid Cartesian product and can get the values according to our requirements.

Joining Tables

- Use a join to query data from more than one table.

```
SELECT      table1.column, table2.column  
FROM        table1, table2  
WHERE       table1.column1 = table2.column2;
```

- Write the join condition in the WHERE clause.
- Prefix the column name with the table name when the same column name appears in more than one table.

What is an Equijoin?

- An equijoin is **a join based on equality or matching column values**. This equality is indicated with an equal sign (=) as the comparison operator in the WHERE clause, as the following query shows.

EMPLOYEES

| EMPLOYEE_ID | DEPARTMENT_ID |
|-------------|---------------|
| 200 | 10 |
| 201 | 20 |
| 202 | 20 |
| 124 | 50 |
| 141 | 50 |
| 142 | 50 |
| 143 | 50 |
| 144 | 50 |
| 103 | 60 |
| 104 | 60 |
| 107 | 60 |
| 149 | 80 |
| 174 | 80 |
| 176 | 80 |

...

DEPARTMENTS

| DEPARTMENT_ID | DEPARTMENT_NAME |
|---------------|-----------------|
| 10 | Administration |
| 20 | Marketing |
| 20 | Marketing |
| 50 | Shipping |
| 60 | IT |
| 60 | IT |
| 60 | IT |
| 80 | Sales |
| 80 | Sales |
| 80 | Sales |

...

Foreign key Primary key

Retrieving Records with Equijoins

```
SELECT employees.employee_id, employees.last_name,  
       employees.department_id, departments.department_id,  
       departments.location_id  
FROM   employees, departments  
WHERE  employees.department_id = departments.department_id;
```

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_ID | LOCATION_ID |
|-------------|-----------|---------------|---------------|-------------|
| 200 | Whalen | 10 | 10 | 1700 |
| 201 | Hartstein | 20 | 20 | 1800 |
| 202 | Fay | 20 | 20 | 1800 |
| 124 | Mourgos | 50 | 50 | 1500 |
| 141 | Rajs | 50 | 50 | 1500 |
| 142 | Davies | 50 | 50 | 1500 |
| 143 | Matos | 50 | 50 | 1500 |
| 144 | Vargas | 50 | 50 | 1500 |
| ... | | | | |

19 rows selected.

Additional Search Conditions Using the AND Operator

```
SELECT last_name, employees.department_id, department_name  
FROM employees, departments  
WHERE employees.department_id = departments.department_id  
AND last_name = 'Matos'
```

EMPLOYEES

| LAST_NAME | DEPARTMENT_ID |
|-----------|---------------|
| Whalen | 10 |
| Hartstein | 20 |
| Fay | 20 |
| Mourgos | 50 |
| Rajs | 50 |
| Davies | 50 |
| Matos | 50 |
| Vargas | 50 |
| Hunold | 60 |
| Ernst | 60 |

...

DEPARTMENTS

| DEPARTMENT_ID | DEPARTMENT_NAME |
|---------------|-----------------|
| 10 | Administration |
| 20 | Marketing |
| 20 | Marketing |
| 50 | Shipping |
| 60 | IT |
| 60 | IT |

...

Qualifying Ambiguous Column Names

- Use table prefixes to qualify column names that are in multiple tables.
- Improve performance by using table prefixes.
- Distinguish columns that have identical names but reside in different tables by using column aliases.

Using Table Aliases

- Simplify queries by using table aliases.
- Improve performance by using table prefixes.

```
SELECT e.employee_id, e.last_name, e.department_id,  
       d.department_id, d.location_id  
  FROM employees e , departments d  
 WHERE e.department_id = d.department_id;
```

Joining More than Two Tables

EMPLOYEES

| LAST_NAME | DEPARTMENT_ID |
|-----------|---------------|
| King | 90 |
| Kochhar | 90 |
| De Haan | 90 |
| Hunold | 60 |
| Ernst | 60 |
| Lorentz | 60 |
| Mourgos | 50 |
| Rajs | 50 |
| Davies | 50 |
| Matos | 50 |
| Vargas | 50 |
| Zlotkey | 80 |
| Abel | 80 |
| Taylor | 80 |
| ... | |

20 rows selected.

DEPARTMENTS

| DEPARTMENT_ID |
|---------------|
| 10 |
| 20 |
| 50 |
| 60 |
| 80 |
| 90 |
| 110 |
| 190 |

8 rows selected.

LOCATIONS

| LOCATION_ID | CITY |
|-------------|---------------------|
| 1400 | Southlake |
| 1500 | South San Francisco |
| 1700 | Seattle |
| 1800 | Toronto |
| 2500 | Oxford |

- To join n tables together, you need a minimum of $n-1$ join conditions. For example, to join three tables, a minimum of two joins is required.

Non-Equijoins

EMPLOYEES

| LAST_NAME | SALARY |
|-----------|--------|
| King | 24000 |
| Kochhar | 17000 |
| De Haan | 17000 |
| Hunold | 9000 |
| Ernst | 6000 |
| Lorentz | 4200 |
| Mourgos | 5800 |
| Rajs | 3500 |
| Davies | 3100 |
| Matos | 2600 |
| Vargas | 2500 |
| Zlotkey | 10500 |
| Abel | 11000 |
| Taylor | 8600 |
| ... | |

20 rows selected.

JOB_GRADES

| GRA | LOWEST_SAL | HIGHEST_SAL |
|-----|------------|-------------|
| A | 1000 | 2999 |
| B | 3000 | 5999 |
| C | 6000 | 9999 |
| D | 10000 | 14999 |
| E | 15000 | 24999 |
| F | 25000 | 40000 |



Salary in the EMPLOYEES table must be between lowest salary and highest salary in the JOB_GRADES table.

Retrieving Records with Non-Equiijoins

```
SELECT e.last_name, e.salary, j.grade_level
FROM employees e, job_grades j
WHERE e.salary
      BETWEEN j.lowest_sal AND j.highest_sal;
```

| LAST_NAME | SALARY | GRA |
|-----------|--------|-----|
| Matos | 2600 | A |
| Vargas | 2500 | A |
| Lorentz | 4200 | B |
| Mourgos | 5800 | B |
| Rajs | 3500 | B |
| Davies | 3100 | B |
| Whalen | 4400 | B |
| Hunold | 9000 | C |
| Ernst | 6000 | C |
| ... | | |

20 rows selected.

Self Joins

EMPLOYEES (WORKER)

| EMPLOYEE_ID | LAST_NAME | MANAGER_ID |
|-------------|-----------|------------|
| 100 | King | |
| 101 | Kochhar | 100 |
| 102 | De Haan | 100 |
| 103 | Hunold | 102 |
| 104 | Ernst | 103 |
| 107 | Lorentz | 103 |
| 124 | Mourgos | 100 |

...

EMPLOYEES (MANAGER)

| EMPLOYEE_ID | LAST_NAME |
|-------------|-----------|
| 100 | King |
| 101 | Kochhar |
| 102 | De Haan |
| 103 | Hunold |
| 104 | Ernst |
| 107 | Lorentz |
| 124 | Mourgos |

...



**MANAGER_ID in the WORKER table is equal to
EMPLOYEE_ID in the MANAGER table.**

Joining a Table to Itself

```
SELECT worker.last_name || ' works for '
    || manager.last_name
FROM   employees worker, employees manager
WHERE  worker.manager_id = manager.employee_id ;
```

| WORKER.LAST_NAME 'WORKSFOR' MANAGER.LAST_NAME |
|---|
| Kochhar works for King |
| De Haan works for King |
| Mourgos works for King |
| Zlotkey works for King |
| Hartstein works for King |
| Whalen works for Kochhar |
| Higgins works for Kochhar |
| Hunold works for De Haan |
| Ernst works for Hunold |
| ... |

19 rows selected.

Creating Cross Joins

- The CROSS JOIN clause produces the cross-product of two tables.
- This is the same as a Cartesian product between the two tables.

```
SELECT last_name, department_name  
FROM employees  
CROSS JOIN departments ;
```

| LAST_NAME | DEPARTMENT_NAME |
|-----------|-----------------|
| King | Administration |
| Kochhar | Administration |
| De Haan | Administration |
| Hunold | Administration |
| ... | |

160 rows selected.

Creating Natural Joins

- The NATURAL JOIN clause is based on all columns in the two tables that have the same name.
- It selects rows from the two tables that have equal values in all matched columns.
- If the columns having the same names have different data types, an error is returned.

Retrieving Records with Natural Joins

```
SELECT department_id, department_name,  
       location_id, city  
FROM   departments  
NATURAL JOIN locations ;
```

| DEPARTMENT_ID | DEPARTMENT_NAME | LOCATION_ID | CITY |
|---------------|-----------------|-------------|---------------------|
| 60 | IT | 1400 | Southlake |
| 50 | Shipping | 1500 | South San Francisco |
| 10 | Administration | 1700 | Seattle |
| 90 | Executive | 1700 | Seattle |
| 110 | Accounting | 1700 | Seattle |
| 190 | Contracting | 1700 | Seattle |
| 20 | Marketing | 1800 | Toronto |
| 80 | Sales | 2500 | Oxford |

8 rows selected.

Creating Joins with the USING Clause

- If several columns have the same names but the data types do not match, the NATURAL JOIN clause can be modified with the **USING clause to specify the columns that should be used for an equijoin.**
- Use the USING clause to match only one column when more than one column matches.
- Do not use a table name or alias in the referenced columns.
- The NATURAL JOIN and USING clauses are mutually exclusive.

Retrieving Records with the USING Clause

```
SELECT e.employee_id, e.last_name, d.location_id  
FROM employees e JOIN departments d  
USING (department_id);
```

| EMPLOYEE_ID | LAST_NAME | LOCATION_ID |
|-------------|-----------|-------------|
| 200 | Whalen | 1700 |
| 201 | Hartstein | 1800 |
| 202 | Fay | 1800 |
| 124 | Mourgos | 1500 |
| 141 | Rajs | 1500 |
| 142 | Davies | 1500 |
| 143 | Matos | 1500 |
| 144 | Vargas | 1500 |
| 103 | Hunold | 1400 |

...
19 rows selected.

```
SELECT employee_id, last_name,  
       employees.department_id, location_id  
FROM   employees, departments  
WHERE  employees.department_id = departments.department_id;
```

Creating Joins with the ON Clause

- The join condition for the natural join is basically an equijoin of all columns with the same name.
- To specify arbitrary conditions or specify columns to join, **the ON clause is used**.
- The join condition is separated from other *search* conditions.
- The ON clause makes code easy to understand.

Retrieving Records with the ON Clause

```
SELECT e.employee_id, e.last_name, e.department_id,  
       d.department_id, d.location_id  
  FROM employees e JOIN departments d  
    ON (e.department_id = d.department_id);
```

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_ID | LOCATION_ID |
|-------------|-----------|---------------|---------------|-------------|
| 200 | Whalen | 10 | 10 | 1700 |
| 201 | Hartstein | 20 | 20 | 1800 |
| 202 | Fay | 20 | 20 | 1800 |
| 124 | Mourgos | 50 | 50 | 1500 |
| 141 | Rajs | 50 | 50 | 1500 |
| 142 | Davies | 50 | 50 | 1500 |
| 143 | Matos | 50 | 50 | 1500 |
| ... | | | | |

19 rows selected.

Creating Three-Way Joins with the ON Clause

```
SELECT employee_id, city, department_name
FROM employees e
JOIN departments d
ON d.department_id = e.department_id
JOIN locations l
ON d.location_id = l.location_id;
```

| EMPLOYEE_ID | CITY | DEPARTMENT_NAME |
|-------------|---------------------|-----------------|
| 103 | Southlake | IT |
| 104 | Southlake | IT |
| 107 | Southlake | IT |
| 124 | South San Francisco | Shipping |
| 141 | South San Francisco | Shipping |
| 142 | South San Francisco | Shipping |
| 143 | South San Francisco | Shipping |
| 144 | South San Francisco | Shipping |
| ... | | |

19 rows selected.

Outer Joins

DEPARTMENTS

| DEPARTMENT_NAME | DEPARTMENT_ID |
|-----------------|---------------|
| Administration | 10 |
| Marketing | 20 |
| Shipping | 50 |
| IT | 60 |
| Sales | 80 |
| Executive | 90 |
| Accounting | 110 |
| Contracting | 190 |

8 rows selected.

EMPLOYEES

| DEPARTMENT_ID | LAST_NAME |
|---------------|-----------|
| 90 | King |
| 90 | Kochhar |
| 90 | De Haan |
| 60 | Hunold |
| 60 | Ernst |
| 60 | Lorentz |
| 50 | Mourgos |
| 50 | Rajs |
| 50 | Davies |
| 50 | Matos |
| 50 | Vargas |
| 80 | Zlotkey |
| ... | |

20 rows selected.

There are no employees in department 190.



Outer Joins Syntax

- You use an outer join to also see rows that do not meet the join condition.
- The left and right joint are the syntax.

```
SELECT      table1.column, table2.column  
FROM        table1 left join table2  
on          table1.column = table2.column;
```

```
SELECT      table1.column, table2.column  
FROM        table1 right join table2  
on          table1.column table2.column;
```

LEFT OUTER JOIN

```
SELECT e.last_name, e.department_id, d.department_name  
FROM employees e  
LEFT OUTER JOIN departments d  
ON (e.department_id = d.department_id) ;
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Whalen | 10 | Administration |
| Fay | 20 | Marketing |
| Hartstein | 20 | Marketing |
| ... | | |
| De Haan | 90 | Executive |
| Kochhar | 90 | Executive |
| King | 90 | Executive |
| Gietz | 110 | Accounting |
| Higgins | 110 | Accounting |
| Grant | | |

20 rows selected.

RIGHT OUTER JOIN

```
SELECT e.last_name, e.department_id, d.department_name  
FROM employees e  
RIGHT OUTER JOIN departments d  
ON (e.department_id = d.department_id) ;
```

| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| King | 90 | Executive |
| Kochhar | 90 | Executive |
| ... | | |
| Whalen | 10 | Administration |
| Hartstein | 20 | Marketing |
| Fay | 20 | Marketing |
| Higgins | 110 | Accounting |
| Gietz | 110 | Accounting |
| | | Contracting |

20 rows selected.

INNER Versus OUTER Joins

- In SQL: 1999, the join of two tables returning only matched rows is an inner join.
- A join between two tables that returns the results of the inner join as well as unmatched rows left (or right) tables is a left (or right) outer join.
- A join between two tables that returns the results of an inner join as well as the results of a left and right join is a full outer join.

FULL OUTER JOIN

```
SELECT e.last_name, e.department_id, d.department_name  
FROM employees e  
FULL OUTER JOIN departments d  
ON (e.department_id = d.department_id) ;
```

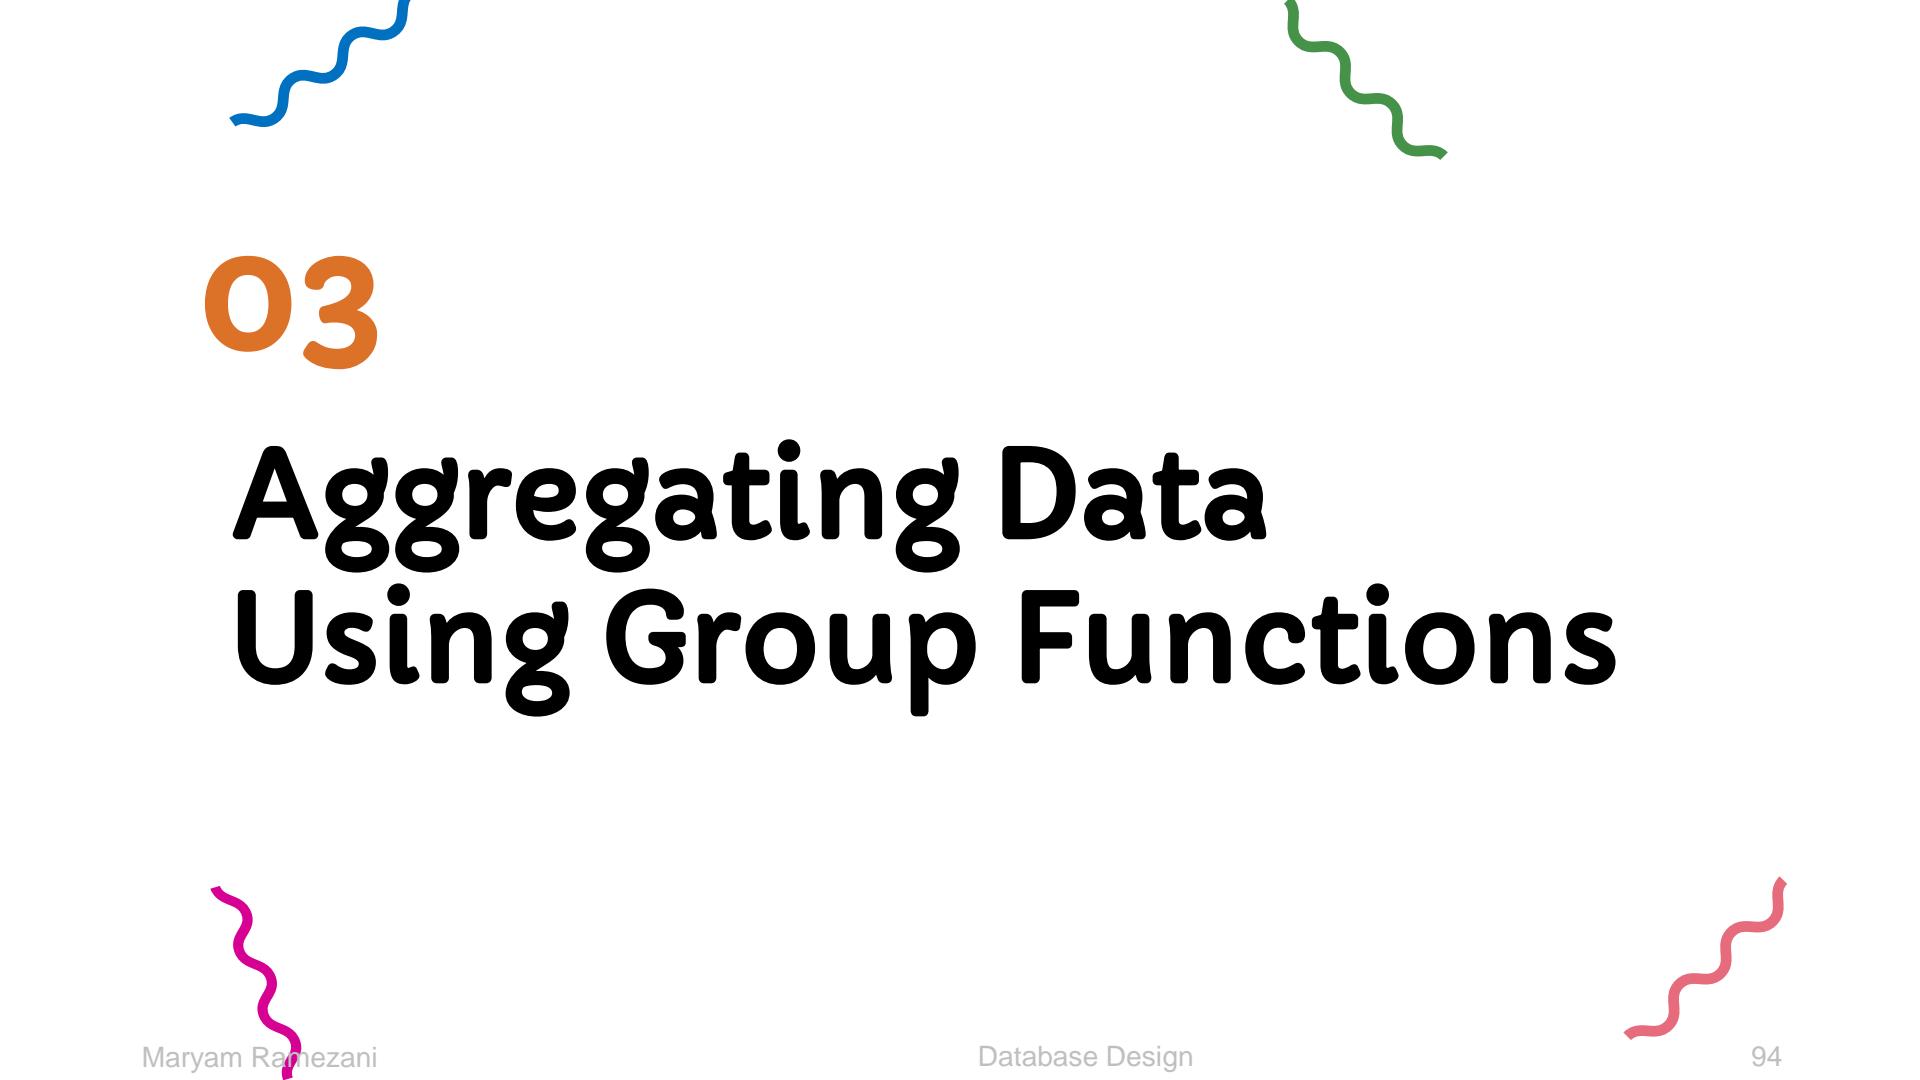
| LAST_NAME | DEPARTMENT_ID | DEPARTMENT_NAME |
|-----------|---------------|-----------------|
| Whalen | 10 | Administration |
| Fay | 20 | Marketing |
| ... | | |
| De Haan | 90 | Executive |
| Kochhar | 90 | Executive |
| King | 90 | Executive |
| Gietz | 110 | Accounting |
| Higgins | 110 | Accounting |
| Grant | | |
| | | Contracting |

21 rows selected.

Additional Conditions

```
SELECT e.employee_id, e.last_name, e.department_id,  
       d.department_id, d.location_id  
FROM   employees e JOIN departments d  
ON     (e.department_id = d.department_id)  
AND    e.manager_id = 149 ;
```

| EMPLOYEE_ID | LAST_NAME | DEPARTMENT_ID | DEPARTMENT_ID | LOCATION_ID |
|-------------|-----------|---------------|---------------|-------------|
| 174 | Abel | | 80 | 80 |
| 176 | Taylor | | 80 | 2500 |



03

Aggregating Data Using Group Functions

What Are Group Functions?

- Group functions operate on sets of rows to give one result per group.

EMPLOYEES

| DEPARTMENT_ID | SALARY |
|---------------|--------|
| 90 | 24000 |
| 90 | 17000 |
| 90 | 17000 |
| 60 | 9000 |
| 60 | 6000 |
| 60 | 4200 |
| 50 | 5800 |
| 50 | 3500 |
| 50 | 3100 |
| 50 | 2600 |
| 50 | 2500 |
| 80 | 10500 |
| 80 | 11000 |
| 80 | 8600 |
| 80 | 7000 |
| 10 | 4400 |

20 rows selected.

The maximum salary in the EMPLOYEES table.

| MAX(SALARY) |
|-------------|
| 24000 |

Types of Group Functions (Aggregations)

- AVG
- COUNT
- MAX
- MIN
- STDDEV
- SUM
- VARIANCE

Group Functions Syntax

```
SELECT [column,] group function(column), ...
      FROM          table
      [WHERE condition]
      [GROUP BY        column]
      [ORDER BY        column];
```

Using the AVG and SUM Functions

- You can use AVG and SUM for numeric data.

```
SELECT AVG(salary), MAX(salary),  
       MIN(salary), SUM(salary)  
FROM   employees  
WHERE  job_id LIKE '%REP%';
```

| AVG(SALARY) | MAX(SALARY) | MIN(SALARY) | SUM(SALARY) |
|-------------|-------------|-------------|-------------|
| 8150 | 11000 | 6000 | 32600 |

Using the MIN and MAX Functions

- You can use MIN and MAX for any data type.

```
SELECT MIN(hire_date), MAX(hire_date)  
FROM employees;
```

| MIN(HIRE_ | MAX(HIRE_ |
|-----------|-----------|
| 17-JUN-87 | 29-JAN-00 |

```
SELECT MIN(last_name), MAX(last_name)  
FROM employees;
```

Using the COUNT Function

- COUNT(*) returns the number of rows in a table.

```
SELECT COUNT(*)  
FROM employees  
WHERE department_id = 50;
```

| COUNT(*) |
|----------|
| 5 |

Using the COUNT Function

- COUNT(expr) returns the number of rows with non-null values for the expr.
- Display the number of department values in the EMPLOYEES table, excluding the null values.

```
SELECT COUNT(commission_pct)
FROM employees
WHERE department_id = 80;
```

| COUNT(COMMISSION_PCT) |
|-----------------------|
| 3 |

Using the DISTINCT Keyword

- COUNT (DISTINCT expr) returns the number of distinct non-null values of the *expr*.
- Display the number of distinct department values in the EMPLOYEES table.

```
SELECT COUNT(DISTINCT department_id)  
FROM employees;
```

| COUNT(DISTINCTDEPARTMENT_ID) |
|------------------------------|
| 7 |

7

Group Functions and Null Values

- Group functions ignore null values in the column.

```
select avg(opening_amt) from customer
```

- The coalesce function forces group functions to include null values.

```
select avg(coalesce(opening_amt,0)) from customer
```

Creating Groups of Data

EMPLOYEES

| DEPARTMENT_ID | SALARY |
|---------------|--------|
| 10 | 4400 |
| 20 | 13000 |
| 20 | 6000 |
| 50 | 5800 |
| 50 | 3500 |
| 50 | 3100 |
| 50 | 2500 |
| 50 | 2600 |
| 60 | 9000 |
| 60 | 6000 |
| 60 | 4200 |
| 80 | 10500 |
| 80 | 8600 |
| 80 | 11000 |
| 90 | 24000 |
| 90 | 17000 |

...
20 rows selected.

4400
9500
3500
6400
10033
The average salary in EMPLOYEES table for each department.
10033

| DEPARTMENT_ID | AVG(SALARY) |
|---------------|-------------|
| 10 | 4400 |
| 20 | 9500 |
| 50 | 3500 |
| 60 | 6400 |
| 80 | 10033.3333 |
| 90 | 19333.3333 |
| 110 | 10150 |
| | 7000 |

Creating Groups of Data: The GROUP BY Clause

Syntax

- Divide rows in a table into smaller groups by using the GROUP BY clause.

```
SELECT column, group_function(column)
FROM          table
[WHERE condition]
[GROUP BY      group_by_expression]
[ORDER BY      column];
```

Using the GROUP BY Clause

- All columns in the SELECT list that are not in group functions must be in the GROUP BY clause.

```
SELECT department_id, AVG(salary)
FROM employees
GROUP BY department_id ;
```

| DEPARTMENT_ID | AVG(SALARY) |
|---------------|-------------|
| 10 | 4400 |
| 20 | 9500 |
| 50 | 3500 |
| 60 | 6400 |
| 80 | 10033.3333 |
| 90 | 19333.3333 |
| 110 | 10150 |
| | 7000 |

8 rows selected.

Using the GROUP BY Clause

- The GROUP BY column does not have to be in the SELECT list.

```
SELECT      AVG(salary)
FROM        employees
GROUP BY    department_id ;
```

| AVG(SALARY) |
|-------------|
| 4400 |
| 9500 |
| 3500 |
| 6400 |
| 10033.3333 |
| 19333.3333 |
| 10150 |
| 7000 |

Grouping by More Than One Column

EMPLOYEES

| DEPARTMENT_ID | JOB_ID | SALARY |
|---------------|------------|--------|
| 90 | AD_PRES | 24000 |
| 90 | AD_VP | 17000 |
| 90 | AD_VP | 17000 |
| 60 | IT_PROG | 9000 |
| 60 | IT_PROG | 6000 |
| 60 | IT_PROG | 4200 |
| 50 | ST_MAN | 5800 |
| 50 | ST_CLERK | 3500 |
| 50 | ST_CLERK | 3100 |
| 50 | ST_CLERK | 2600 |
| 50 | ST_CLERK | 2500 |
| 80 | SA_MAN | 10500 |
| 80 | SA_REP | 11000 |
| 80 | SA_REP | 8600 |
| ... | | |
| 20 | MK_REP | 6000 |
| 110 | AC_MGR | 12000 |
| 110 | AC_ACCOUNT | 8300 |

20 rows selected.

“Add up the salaries in the EMPLOYEES table for each job, grouped by department.”

| DEPARTMENT_ID | JOB_ID | SUM(SALARY) |
|---------------|------------|-------------|
| 10 | AD_ASST | 4400 |
| 20 | MK_MAN | 13000 |
| 20 | MK_REP | 6000 |
| 50 | ST_CLERK | 11700 |
| 50 | ST_MAN | 5800 |
| 60 | IT_PROG | 19200 |
| 80 | SA_MAN | 10500 |
| 80 | SA_REP | 19600 |
| 90 | AD_PRES | 24000 |
| 90 | AD_VP | 34000 |
| 110 | AC_ACCOUNT | 8300 |
| 110 | AC_MGR | 12000 |
| | SA_REP | 7000 |

13 rows selected.

Using the GROUP BY Clause on Multiple Columns

```
SELECT department_id dept_id, job_id, SUM(salary)
FROM employees
GROUP BY department_id, job_id ;
```

| DEPT_ID | JOB_ID | SUM(SALARY) |
|---------|------------|-------------|
| 10 | AD_ASST | 4400 |
| 20 | MK_MAN | 13000 |
| 20 | MK_REP | 6000 |
| 50 | ST_CLERK | 11700 |
| 50 | ST_MAN | 5800 |
| 60 | IT_PROG | 19200 |
| 80 | SA_MAN | 10500 |
| 80 | SA_REP | 19600 |
| 90 | AD_PRES | 24000 |
| 90 | AD_VP | 34000 |
| 110 | AC_ACCOUNT | 8300 |
| 110 | AC_MGR | 12000 |
| | SA_REP | 7000 |

13 rows selected.

Illegal Queries - Using Group Functions

- Any column or expression in the SELECT list that is not an aggregate function must be in the GROUP BY clause.
- Column missing in the GROUP BY clause:

```
SELECT department_id, COUNT(last_name)
  FROM employees;
```

```
SELECT department_id, COUNT(last_name)
      *
ERROR at line 1:
ORA-00937: not a single-group group function
```

Illegal Queries - Using Group Functions

- You cannot use the WHERE clause to restrict groups.
- You use the HAVING clause to restrict groups.
- You cannot use group functions in the WHERE clause.
- Cannot use the WHERE clause to restrict groups:

```
SELECT      department_id, AVG(salary)
FROM        employees
WHERE       AVG(salary) > 8000
GROUP BY    department_id;
```

```
WHERE      AVG(salary) > 8000
*
ERROR at line 3:
ORA-00934: group function is not allowed here
```

Excluding Group Results

EMPLOYEES

| DEPARTMENT_ID | SALARY |
|---------------|--------|
| 90 | 24000 |
| 90 | 17000 |
| 90 | 17000 |
| 60 | 9000 |
| 60 | 6000 |
| 60 | 4200 |
| 50 | 5800 |
| 50 | 3500 |
| 50 | 3100 |
| 50 | 2600 |
| 50 | 2500 |
| 80 | 10500 |
| 80 | 11000 |
| 80 | 8600 |
| ... | |
| 20 | 6000 |
| 110 | 12000 |
| 110 | 8300 |

20 rows selected.

The maximum salary per department when it is greater than \$10,000

| DEPARTMENT_ID | MAX(SALARY) |
|---------------|-------------|
| 20 | 13000 |
| 80 | 11000 |
| 90 | 24000 |
| 110 | 12000 |

Excluding Group Results: The HAVING Clause

Use the HAVING clause to restrict groups:

1. Rows are grouped.
2. The group function is applied.
3. Groups matching the HAVING clause are displayed.

```
SELECT column, group_function
      FROM          table
      [WHERE condition]
      [GROUP BY      group_by_expression]
      [HAVING        group_condition]
      [ORDER BY      column] ;
```

Using the HAVING Clause

```
SELECT      department_id, MAX(salary)
FROM        employees
GROUP BY    department_id
HAVING      MAX(salary)>10000 ;
```

| DEPARTMENT_ID | MAX(SALARY) |
|---------------|-------------|
| 20 | 13000 |
| 80 | 11000 |
| 90 | 24000 |
| 110 | 12000 |

Using the HAVING Clause

```
SELECT      job_id, SUM(salary) PAYROLL
FROM        employees
WHERE       job_id NOT LIKE '%REP%'
GROUP BY   job_id
HAVING     SUM(salary) > 13000
ORDER BY   SUM(salary);
```

| JOB_ID | PAYROLL |
|---------|---------|
| IT_PROG | 19200 |
| AD_PRES | 24000 |
| AD_VP | 34000 |

Nesting Group Functions

- Display the maximum average salary.

```
SELECT MAX(AVG(salary))  
FROM employees  
GROUP BY department_id;
```

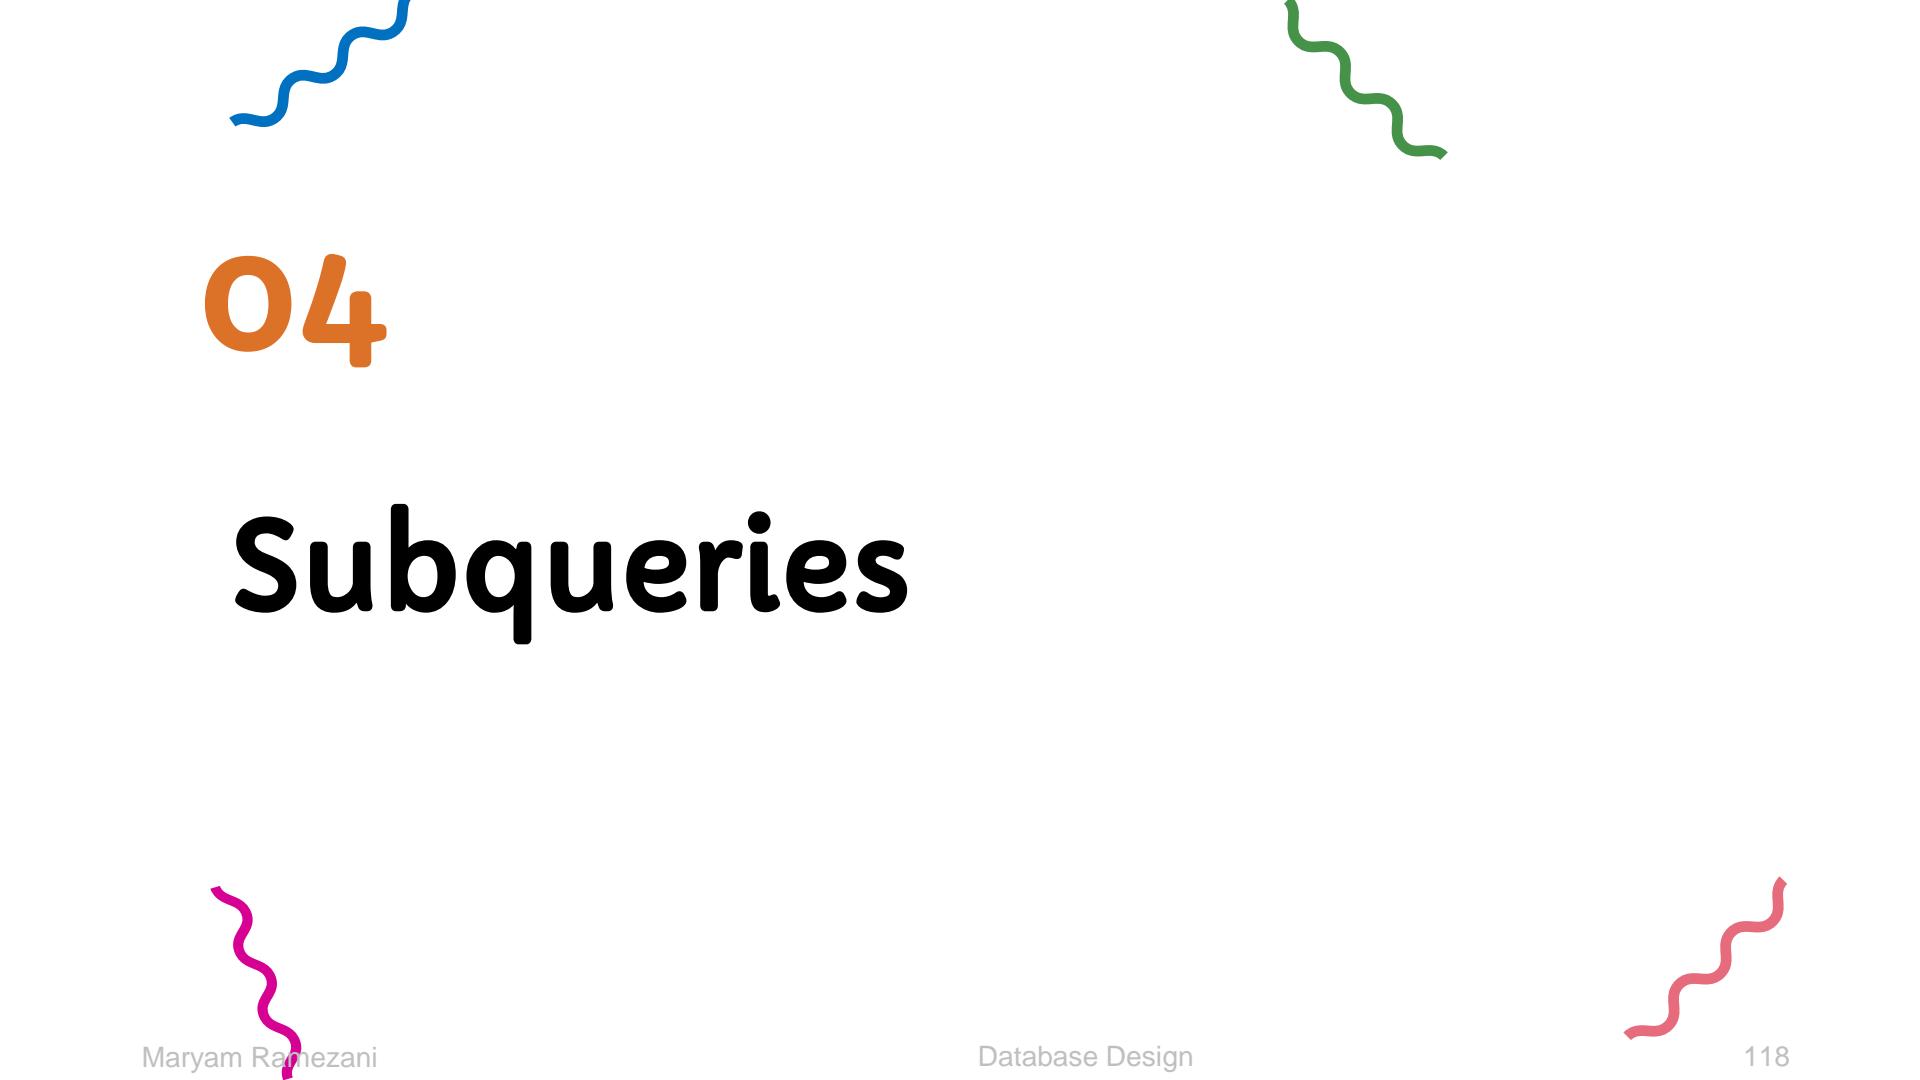
| MAX(AVG(SALARY)) |
|------------------|
| 19333.3333 |

Summary

You should have learned how to:

- Use the group functions COUNT, MAX, MIN, AVG
- Write queries that use the GROUP BY clause
- Write queries that use the HAVING clause

```
SELECT column, group_function(column)
FROM           table
[WHERE condition]
[GROUP BY      group_by_expression]
[HAVING        group_condition]
[ORDER BY      column];
```



04

Subqueries

Using a Subquery to Solve a Problem

Who has a salary greater than Abel's?

Main Query:



Which employees have salaries greater than Abel's salary?



Subquery



What is Abel's salary?

Nested Queries - Definitions

- A nested query is a query inside another query
 - The enclosing query also called outer query.
 - Nested query is called inner query.
- It usually appears as a condition in where or having clauses.
- There can be multiple levels of nesting
- There are two kinds of nested queries
 - Correlated
 - Non-Correlated

Example:

```
Select movie_title  
From movies  
Where director_id IN (  
    Select person_id  
    From People  
    Where person_state = 'TX')
```

Nested Queries: Non-Correlated

- Generates data required by outer query before it can be executed
- Inner query does not contain any reference to outer query
- Behaves like a procedure
- The result should not contain any column from the nested query
- Example:

Schema:

- People(person_fname, person_lname, person_id, person_state, person_city)
- Movies(movie_id, movie_title, director_id, studio_id)

Query: Select movie_title, studio_id
From Movies
Where director_id IN (Select person_id
From People
Where person_state = 'TX')

Steps:

- Subquery is executed
- Subquery results are plugged into the outer query
- The outer query is processed

Nested Queries: Correlated

- Contains reference to the outer query.
- Behaves like a loop.

Example:

Schema: People(person_fname, person_lname, person_id, person_state, person_city)
 Cast_Movies(cast_member_id, role, movie_id)

Query: select person_fname, person_lname
 From People p1
 Where 'Pam Green' in (Select role
 From Cast_Movies
 Where p1.person_id = cast_member_id)

Steps:

- Contents of the table row in outer query are read
- Sub-query is executed using data in the row being processed.
- Results of the inner query are passed to the where in the outer query
- The Outer query is Processed
- Loop continues till all rows are exhausted

Subquery Syntax

- The subquery (inner query) executes once before the main query.
- The result of the subquery is used by the main query (outer query).

```
SELECT      select_list
FROM        table
WHERE       expr operator
            (SELECT      select_list
             FROM table);
```

Using a Subquery

```
SELECT last_name
FROM   employees
WHERE  salary > 11000
       (SELECT salary
        FROM   employees
        WHERE  last_name = 'Abel');
```

| LAST_NAME |
|-----------|
| King |
| Kochhar |
| De Haan |
| Hartstein |
| Higgins |

Guidelines for Using Subqueries

- Enclose subqueries in parentheses.
- Place subqueries on the right side of the comparison condition.
- The ORDER BY clause in the subquery is not needed unless you are performing Top-N analysis.
- Use single-row operators with single-row subqueries and use multiple-row operators with multiple-row subqueries.

Types of Subqueries

- Types of Subqueries
 - **Single-row subqueries:** Queries that return only one row from the inner SELECT statement
 - **Multiple-row subqueries:** Queries that return more than one row from the inner SELECT statement
 - **Note:** There are also **multiple-column subqueries:** Queries that return more than one column from the inner SELECT statement
 - Single-row subquery



- Multiple-row subquery



Single-Row Subqueries

- Return only one row
- Use single-row comparison operators

| Operator | Meaning |
|----------|---------------------------------|
| = | Equal to |
| > | Greater than |
| >= | Greater than or equal to |
| < | Less than |
| <= | Less than or equal to |
| <> | Not equal to |

Executing Single-Row Subqueries

```
SELECT last_name, job_id, salary
FROM employees
WHERE job_id = ST_CLERK
      (SELECT job_id
       FROM employees
       WHERE employee_id = 141)
AND salary > 2600
      (SELECT salary
       FROM employees
       WHERE employee_id = 143);
```

| LAST_NAME | JOB_ID | SALARY |
|-----------|----------|--------|
| Rajs | ST_CLERK | 3500 |
| Davies | ST_CLERK | 3100 |

Using Group Functions in a Subquery

```
SELECT last_name, job_id, salary  
FROM employees  
WHERE salary =  
      (SELECT MIN(salary)  
       FROM employees);
```

| LAST_NAME | JOB_ID | SALARY |
|-----------|----------|--------|
| Vargas | ST_CLERK | 2500 |

The HAVING Clause with Subqueries

- The Oracle server executes subqueries first.
- The Oracle server returns results into the HAVING clause of the main query.

```
SELECT      department_id, MIN(salary)
FROM        employees
GROUP BY    department_id
HAVING      MIN(salary) > 2500
            (SELECT MIN(salary)
             FROM   employees
             WHERE  department_id = 50);
```

What is Wrong with this Statement?

```
SELECT employee_id, last_name  
FROM   employees  
WHERE  salary =  
       (SELECT    MIN(salary)  
        FROM      employees  
        GROUP BY department_id);
```

```
ERROR at line 4:  
ORA-01427: single-row subquery returns more than  
one row
```

Single-row operator with multiple-row
subquery

Will this Statement Return Rows?

```
SELECT last_name, job_id  
FROM employees  
WHERE job_id =  
      (SELECT job_id  
       FROM employees  
       WHERE last_name = 'Haas');
```

```
no rows selected
```

Subquery returns no values

Multiple-Row Subqueries

- Return more than one row
- Use multiple-row comparison operators

| Operator | Meaning |
|----------|---|
| IN | Equal to any member in the list |
| ANY | Compare value to each value returned by the subquery |
| ALL | Compare value to every value returned by the subquery |

Using the ANY Operator in Multiple-Row Subqueries

```
SELECT employee_id, last_name, job_id, salary
FROM   employees      9000, 6000, 4200
WHERE  salary < ANY
       (SELECT salary
        FROM   employees
        WHERE  job_id = 'IT_PROG')
AND    job_id <> 'IT_PROG';
```

| EMPLOYEE_ID | LAST_NAME | JOB_ID | SALARY |
|-------------|-----------|----------|--------|
| 124 | Mourgos | ST_MAN | 5800 |
| 141 | Rajs | ST_CLERK | 3500 |
| 142 | Davies | ST_CLERK | 3100 |
| 143 | Matos | ST_CLERK | 2600 |
| 144 | Vargas | ST_CLERK | 2500 |
| ... | | | |

10 rows selected.

Using the ALL Operator in Multiple-Row Subqueries

```
SELECT employee_id, last_name, job_id, salary
FROM   employees
WHERE  salary < ALL
       (SELECT salary
        FROM   employees
        WHERE  job_id = 'IT_PROG')
AND    job_id <> 'IT_PROG';
```

9000, 6000, 4200

| EMPLOYEE_ID | LAST_NAME | JOB_ID | SALARY |
|-------------|-----------|----------|--------|
| 141 | Rajs | ST_CLERK | 3500 |
| 142 | Davies | ST_CLERK | 3100 |
| 143 | Matos | ST_CLERK | 2600 |
| 144 | Vargas | ST_CLERK | 2500 |

Null Values in a Subquery

```
SELECT emp.last_name
FROM   employees emp
WHERE  emp.employee_id NOT IN
       (SELECT mgr.manager_id
        FROM   employees mgr);
no rows selected
```

Any, Some, All

$$\left\{ \begin{array}{ll} \text{theta} & \text{ANY} \\ \text{theta} & \text{SOME} \\ \text{theta} & \text{ALL} \end{array} \right\}$$
$$\text{theta} \in \left\{ \begin{array}{l} = \\ \neq \\ < \\ \leq \\ \geq \\ > \end{array} \right\}$$

- The ALL keyword modifies the greater than comparison operator to mean greater than all values.
- The ANY keyword is not as restrictive as the ALL keyword.
- When used with the greater than comparison operator, "> ANY" means greater than some value.
- The "= ANY" operator is exactly equivalent to the IN operator.
- However, the "!= ANY" (not equal any) is not equivalent to the NOT IN operator.

Any, Some, All

- Give the providers whose status are not maximum.

1- SELECT S#

FROM S

WHERE STATUS < ANY (SELECT DISTINCT STATUS FROM S)

2- SELECT S#

FROM S

WHERE STATUS < (SELECT MAX (STATUS) FROM S)

Subqueries and the EXISTS Operator

- ❑ When a subquery uses the EXISTS operator, the subquery functions as an existence test.
- ❑ The WHERE clause of the outer query tests for the existence of rows returned by the inner query.
- ❑ The subquery does not actually produce any data; rather, it returns a value of TRUE or FALSE.
- ❑ The general format of a subquery WHERE clause with an EXISTS operator is shown here.
- ❑ Note that the NOT operator can also be used to negate the result of the EXISTS operator.

WHERE [NOT] EXISTS (subquery)

Example

```
SELECT emp_last_name "Last Name", emp_first_name "First Name"  
FROM employee  
WHERE EXISTS  
  (SELECT *  
   FROM dependent  
   WHERE emp_ssn = dep_emp_ssn);
```

| Last Name | First Name |
|-----------|------------|
| Joyner | Suzanne |
| Zhu | Waiman |
| Bock | Douglas |

Subqueries and the EXISTS operator

- ❑ Subqueries using an EXISTS operator are a bit different from other subqueries, in the following ways:
 - The keyword EXISTS is not preceded by a column name, constant, or other expression.
 - The SELECT clause list of a subquery that uses an EXISTS operator almost always consists of an asterisk (*). This is because there is no real point in listing column names since you are simply testing for the existence of rows that meet the conditions specified in the subquery.
 - The subquery evaluates to TRUE or FALSE rather than returning any data.
 - A subquery that uses an EXISTS operator will always be a correlated subquery.

Subqueries and the EXISTS operator

- ❑ The EXISTS operator is very important, because there is often no alternative to its use.
- ❑ All queries that use the IN operator or a modified comparison operator (=, <, >, etc. modified by ANY or ALL) can be expressed with the EXISTS operator.
- ❑ However, some queries formulated with EXISTS cannot be expressed in any other way!
- ❑ The NOT EXISTS operator is the mirror-image of the EXISTS operator.
- ❑ A query that uses NOT EXISTS in the WHERE clause is satisfied if the subquery returns no rows.

Subqueries and the EXISTS operator

```
SELECT emp_last_name  
FROM employee  
WHERE emp_ssn = ANY  
  (SELECT dep_emp_ssn  
   FROM dependent);
```

| EMP_LAST_NAME |
|---------------|
| ----- |
| Bock |
| Zhu |
| Joyner |

```
SELECT  
      emp_last_name  
  FROM employee  
 WHERE EXISTS  
       (SELECT *  
        FROM dependent  
       WHERE emp_ssn  
         = dep_emp_ssn);
```

| EMP_LAST_NAME |
|---------------|
| ----- |
| Bock |
| Zhu |
| Joyner |

Subqueries and the ORDER BY Clause

- The SELECT statement shown below adds the ORDER BY clause to specify sorting by first name within last name.
- Note that the ORDER BY clause is placed after the WHERE clause, and that this includes the subquery as part of the WHERE clause.

```
SELECT emp_last_name "Last Name",
       emp_first_name "First Name"
  FROM employee
 WHERE EXISTS
       (SELECT *
          FROM dependent
         WHERE emp_ssn = dep_emp_ssn)
 ORDER BY emp_last_name, emp_first_name;
```

Output:

| Last Name | First Name |
|-----------|------------|
| Bock | Douglas |
| Joyner | Suzanne |
| Zhu | Waiman |

Union

- Union Joins allow multiple query results to be combined into a single result set

Syntax

```
Select select_list  
From table [,table, ....]  
[Where condition]  
Union [All]  
Select select_list  
From table [,table, ....]  
[Where condition]
```

Example

```
Select person_id,  
person_city, person_state  
From People  
Union  
Select studio_id,  
studio_city,  
studio_state  
From Studios
```

- Notes:
 - The number of columns selected for both the queries should be the same
 - The columns are merged in order in which they are selected
 - **The duplicates are eliminated from the combined table**
 - More than two tables can be joined together

Union (All & Order By)

- Union query eliminates all duplicates in the resultant table
 - Union All is used when we do not want to eliminate the duplicates
 - Union and Union distinct are the same.
- Union and Order By can be used together to order the results of the combined table
 - This clause is not allowed when a single column result is obtained and the all keyword is used since the duplicates are eliminated and there is nothing to order by
- Example

```
Select studio_id, studio_state  
From Studios  
Union  
Select Person_id, person_state  
From People  
Order By studio_state
```

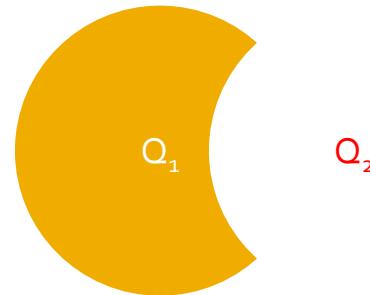
Intersect

- In the Intersect Query results of two separate queries are concatenated, however, only common elements of the two queries are included in the resultset
- Example**

```
Select person_state  
From People  
Intersect  
Select studio_state  
From Studios
```

Except

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
EXCEPT  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```



Subquery Benefits

- They can simplify the logic and readability of your query, especially if you need to filter or aggregate data before joining it with another table.
- They can help you avoid duplicate rows or columns that might result from a join operation.
- They can enable you to perform complex calculations or comparisons that might not be possible with a join.
 - For example, you can use a subquery to find the average salary of each department, and then compare it with the salary of each employee in the main query.

Subquery Drawbacks

Subqueries also have some drawbacks that can affect database performance.

- They can increase the processing time and memory usage of your query, especially if the subquery returns a large number of rows or columns.
- They can limit the optimization options of the database system, as some subqueries cannot use indexes or other techniques to speed up the execution.
- They can introduce errors or inconsistencies if the subquery is not correlated with the main query, or if the subquery data changes during the execution of the main query.

Join Benefits

Joins are another way to query data from multiple tables in a database.

- They can reduce the number of queries and subqueries needed to retrieve the data you want, which can save processing time and memory.
- They can leverage the indexes and other features of the database system to optimize the join operation and make it faster and more efficient.
- They can ensure the consistency and accuracy of the data, as the join condition determines which rows from each table are matched and returned.

Join Drawbacks

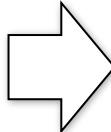
- They can complicate the syntax and readability of your query, especially if you need to join multiple tables or use different types of joins.
- They can generate unwanted or redundant rows or columns that might affect the quality and size of the result set.
- They can require careful planning and design of the database schema and the join condition, as poorly structured or indexed tables or columns can slow down or fail the join operation.

How to choose

- Deciding whether to use a subquery or a join for your query is dependent on various factors, such as the data structure, the query complexity, the database system, and the performance goals.
- As a **general guideline**, you should use a **subquery if you need to filter or aggregate data before joining it with another table**, or **if you need to perform calculations or comparisons that are not possible with a join**. On the other hand, if you need to query data from multiple tables based on a common column or condition, or if you want to take advantage of the optimization features of the database system, then using a join is recommended.
- Ultimately, it is best to test and compare the execution time and result set of both options and choose the one that meets your requirements and expectations.

Nested queries as alternatives to INTERSECT and EXCEPT

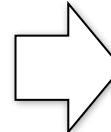
```
(SELECT R.A, R.B  
FROM R)  
INTERSECT  
(SELECT S.A, S.B  
FROM S)
```



```
SELECT R.A, R.B  
FROM R  
WHERE EXISTS (  
    SELECT *  
    FROM S  
    WHERE R.A=S.A  
    AND R.B=S.B)
```

INTERSECT and EXCEPT
not in some DBMSs!

```
(SELECT R.A, R.B  
FROM R)  
EXCEPT  
(SELECT S.A, S.B  
FROM S)
```



```
SELECT R.A, R.B  
FROM R  
WHERE NOT EXISTS (  
    SELECT *  
    FROM S  
    WHERE R.A=S.A AND  
    R.B=S.B)
```

If R, S have no duplicates, then
can write without sub-queries
(HOW?)

Manipulating Data

Adding a New Row to a Table

DEPARTMENTS

| DEPARTMENT_ID | DEPARTMENT_NAME | MANAGER_ID | LOCATION_ID |
|---------------|-----------------|------------|-------------|
| 10 | Administration | 200 | 1700 |
| 20 | Marketing | 201 | 1800 |
| 50 | Shipping | 124 | 1500 |
| 60 | IT | 103 | 1400 |
| 80 | Sales | 149 | 2500 |
| 90 | Executive | 100 | 1700 |
| 110 | Accounting | 205 | 1700 |
| 190 | Contracting | | 1700 |

70 Public Relations

100

1700

New
row

...insert a new row
into the
DEPARTMENTS
table...



| DEPARTMENT_ID | DEPARTMENT_NAME | MANAGER_ID | LOCATION_ID |
|---------------|-----------------|------------|-------------|
| 10 | Administration | 200 | 1700 |
| 20 | Marketing | 201 | 1800 |
| 50 | Shipping | 124 | 1500 |
| 60 | IT | 103 | 1400 |
| 80 | Sales | 149 | 2500 |
| 90 | Executive | 100 | 1700 |
| 110 | Accounting | 205 | 1700 |
| 190 | Contracting | | 1700 |

| | | | |
|----|------------------|-----|------|
| 70 | Public Relations | 100 | 1700 |
|----|------------------|-----|------|

The INSERT Statement Syntax

- Add new rows to a table by using the INSERT statement.
- Only one row is inserted at a time with this syntax.

```
INSERT INTO table [(column [, column...])]  
VALUES  
      (value [, value...]);
```

Inserting New Rows

- Insert a new row containing values for each column.
- List values in the default order of the columns in the table.
- Optionally, list the columns in the INSERT clause.

```
INSERT INTO departments(department_id, department_name,  
                      manager_id, location_id)  
VALUES      (70, 'Public Relations', 100, 1700);  
1 row created.
```

- Enclose character and date values within single quotation marks.

Inserting Rows with Null Values

- Implicit method: Omit the column from the column list.

```
INSERT INTO departments (department_id,  
                      department_name) [] []  
VALUES (30, 'Purchasing');  
1 row created.
```

- **Explicit method: Specify the `NULL` keyword in the `VALUES` clause.**

```
INSERT INTO departments  
VALUES (100, 'Finance', NULL, NULL);  
1 row created.
```

Inserting Special Values

- The `current_date` function records the current date.

```
INSERT INTO employees (employee_id,
                      first_name, last_name,
                      email, phone_number,
                      hire_date, job_id, salary,
                      commission_pct, manager_id,
                      department_id)
VALUES          (113,
                  'Louis', 'Popp',
                  'LPOPP', '515.124.4567',
                  current_date, 'AC_ACCOUNT', 6900,
                  NULL, 205, 100);

1 row created.
```

Creating a Script

- Use : substitution in a SQL statement to prompt for values.
- : is a placeholder for the variable value.

```
INSERT INTO departments
    (department_id, department_name, location_id)
VALUES (:department_id, ':department_name', :location);
```

Define Substitution Variables

| | |
|-------------------|--|
| "department_id" | <input type="text" value="40"/> |
| "department_name" | <input type="text" value="Human Resources"/> |
| "location" | <input type="text" value="2500"/> |

Submit for Execution

Cancel

1 row created.

Copying Rows from Another Table

- Write your INSERT statement with a subquery.

```
INSERT INTO sales_reps(id, name, salary, commission_pct)
    SELECT employee_id, last_name, salary, commission_pct
        FROM employees
    WHERE job_id LIKE '%REP%';
```

4 rows created.

- Do not use the VALUES clause.
- Match the number of columns in the INSERT clause to those in the subquery.

Changing Data in a Table

EMPLOYEES

| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | EMAIL | HIRE_DATE | JOB_ID | SALARY | DEPARTMENT_ID | COMMISSION_PCT |
|-------------|------------|-----------|----------|-----------|---------|--------|---------------|----------------|
| 100 | Steven | King | SKING | 17-JUN-87 | AD_PRES | 24000 | 90 | |
| 101 | Neena | Kochhar | NKOCHHAR | 21-SEP-89 | AD_VP | 17000 | 90 | |
| 102 | Lex | De Haan | LDEHAAN | 13-JAN-93 | AD_VP | 17000 | 90 | |
| 103 | Alexander | Hunold | AHUNOLD | 03-JAN-90 | IT_PROG | 9000 | 60 | |
| 104 | Bruce | Ernst | BERNST | 21-MAY-91 | IT_PROG | 6000 | 60 | |
| 107 | Diana | Lorentz | DLORENTZ | 07-FEB-99 | IT_PROG | 4200 | 60 | |
| 124 | Kevin | Mourgos | KMOURGOS | 16-NOV-99 | ST_MAN | 5800 | 50 | |

Update rows in the EMPLOYEES table.



| EMPLOYEE_ID | FIRST_NAME | LAST_NAME | EMAIL | HIRE_DATE | JOB_ID | SALARY | DEPARTMENT_ID | COMMISSION_PCT |
|-------------|------------|-----------|----------|-----------|---------|--------|---------------|----------------|
| 100 | Steven | King | SKING | 17-JUN-87 | AD_PRES | 24000 | 90 | |
| 101 | Neena | Kochhar | NKOCHHAR | 21-SEP-89 | AD_VP | 17000 | 90 | |
| 102 | Lex | De Haan | LDEHAAN | 13-JAN-93 | AD_VP | 17000 | 90 | |
| 103 | Alexander | Hunold | AHUNOLD | 03-JAN-90 | IT_PROG | 9000 | 30 | |
| 104 | Bruce | Ernst | BERNST | 21-MAY-91 | IT_PROG | 6000 | 30 | |
| 107 | Diana | Lorentz | DLORENTZ | 07-FEB-99 | IT_PROG | 4200 | 30 | |
| 124 | Kevin | Mourgos | KMOURGOS | 16-NOV-99 | ST_MAN | 5800 | 50 | |

The UPDATE Statement Syntax

- Modify existing rows with the UPDATE statement.

```
UPDATE      table
SET         column = value [, column = value, ...]
[WHERE      condition];
```

- Update more than one row at a time, if required.

Updating Rows in a Table

- Specific row or rows are modified if you specify the WHERE clause.

```
UPDATE employees  
SET department_id = 70  
WHERE employee_id = 113;  
1 row updated.
```

- All rows in the table are modified if you omit the WHERE clause.

```
UPDATE copy_emp  
SET department_id = 110;  
22 rows updated.
```

Updating Two Columns with a Subquery

- Update employee 114's job and salary to match that of employee 205.

```
UPDATE      employees
SET          job_id   =  (SELECT    job_id
                         FROM      employees
                         WHERE     employee_id = 205),
            salary    =  (SELECT    salary
                         FROM      employees
                         WHERE     employee_id = 205)
WHERE        employee_id    =  114;
1 row updated.
```

Updating Rows Based on Another Table

- Use subqueries in UPDATE statements to update rows in a table based on values from another table.

```
UPDATE copy_emp
SET department_id = (SELECT department_id
                      FROM employees
                      WHERE employee_id = 100)
WHERE job_id          = (SELECT job_id
                         FROM employees
                         WHERE employee_id = 200);
1 row updated.
```

Updating Rows: Integrity Constraint Error

Department number 55 does not exist in the parent table!

```
UPDATE employees  
SET department_id = 55  
WHERE department_id = 110;
```

```
UPDATE employees  
*  
ERROR at line 1:  
ORA-02291: integrity constraint (HR.EMP_DEPT_FK)  
violated - parent key not found
```

Removing a Row from a Table

DEPARTMENTS

| DEPARTMENT_ID | DEPARTMENT_NAME | MANAGER_ID | LOCATION_ID |
|---------------|-----------------|------------|-------------|
| 10 | Administration | 200 | 1700 |
| 20 | Marketing | 201 | 1800 |
| 30 | Purchasing | | |
| 100 | Finance | | |
| 50 | Shipping | 124 | 1500 |
| 60 | IT | 103 | 1400 |

Delete a row from the DEPARTMENTS table.

| DEPARTMENT_ID | DEPARTMENT_NAME | MANAGER_ID | LOCATION_ID |
|---------------|-----------------|------------|-------------|
| 10 | Administration | 200 | 1700 |
| 20 | Marketing | 201 | 1800 |
| 30 | Purchasing | | |
| 50 | Shipping | 124 | 1500 |
| 60 | IT | 103 | 1400 |

The DELETE Statement

- You can remove existing rows from a table by using the DELETE statement.

```
DELETE [FROM]          table  
[WHERE   condition];
```

Deleting Rows from a Table

- Specific rows are deleted if you specify the WHERE clause.

```
DELETE FROM departments  
WHERE department_name = 'Finance';  
1 row deleted.
```

- All rows in the table are deleted if you omit the WHERE clause.

```
DELETE FROM copy_emp;  
22 rows deleted.
```

Deleting Rows Based on Another Table

- Use subqueries in `DELETE` statements to remove rows from a table based on values from another table.

```
DELETE FROM employees
WHERE department_id =
      (SELECT department_id
       FROM departments
       WHERE department_name LIKE '%Public%');

1 row deleted.
```

Deleting Rows: Integrity Constraint Error

You cannot delete a row that contains a primary key that is used as a foreign key in another table.

```
DELETE FROM departments  
WHERE      department_id = 60;
```

```
DELETE FROM departments  
*  
ERROR at line 1:  
ORA-02292: integrity constraint (HR.EMP_DEPT_FK)  
violated - child record found
```

Overview of the Explicit Default Feature

- With the explicit default feature, you can use the DEFAULT keyword as a column value where the column default is desired.
- The addition of this feature is for compliance with the SQL: 1999 Standard.
- This allows the user to control where and when the default value should be applied to data.
- Explicit defaults can be used in INSERT and UPDATE statements.

Using Explicit Default Values

- DEFAULT with INSERT:

```
INSERT INTO departments
(department_id, department_name, manager_id)
VALUES (300, 'Engineering', DEFAULT);
```

- DEFAULT with UPDATE:

```
UPDATE departments
SET manager_id = DEFAULT WHERE department_id = 10;
```

- If no default value for the corresponding column has been specified, Postgres sets the column to null.

The MERGE Statement

- Provides the ability to **conditionally update or insert data** into a database table
- **Performs an UPDATE if the row exists, and an INSERT if it is a new row:**
 - Avoids separate updates
 - Increases performance and ease of use
 - Is useful in data warehousing applications: you may need to work with data coming from multiple sources, some of which may be duplicates. With the MERGE statement, you can conditionally add or modify rows.
- The MERGE statement is deterministic. You cannot update the same row of the target table multiple times in the same MERGE statement.

The MERGE Statement Syntax

- You can conditionally insert or update rows in a table by using the MERGE statement.

```
MERGE INTO table_name table_alias
    USING (table|view|sub_query) alias
    ON (join condition)
    WHEN MATCHED THEN
        UPDATE SET
            col1 = col_val1,
            col2 = col2_val
    WHEN NOT MATCHED THEN
        INSERT (column_list)
        VALUES (column_values);
```

Merging Rows

- Insert or update rows in the COPY_EMP table to match the EMPLOYEES table.
 - The example shown matches the EMPLOYEE_ID in the COPY_EMP table to the EMPLOYEE_ID in the EMPLOYEES table. If a match is found, the row in the COPY_EMP table is updated to match the row in the EMPLOYEES table. If the row is not found, it is inserted into the COPY_EMP table.

```
MERGE INTO copy_emp c
  USING employees e
  ON (c.employee_id = e.employee_id)
WHEN MATCHED THEN
  UPDATE SET
    c.first_name      = e.first_name,
    c.last_name       = e.last_name,
    ...
    c.department_id  = e.department_id
WHEN NOT MATCHED THEN
  INSERT VALUES(e.employee_id, e.first_name, e.last_name,
                e.email, e.phone_number, e.hire_date, e.job_id,
                e.salary, e.commission_pct, e.manager_id,
                e.department_id);
```

Merging Rows

- The condition `c.employee_id = e.employee_id` is evaluated. Because the `COPY_EMP` table is empty, the condition returns false: there are no matches. The logic falls into the `WHEN NOT MATCHED` clause, and the `MERGE` command inserts the rows of the `EMPLOYEES` table into the `COPY_EMP` table.
- If rows existed in the `COPY_EMP` table and employee IDs matched in both tables (the `COPY_EMP` and `EMPLOYEES` tables), the existing rows in the `COPY_EMP` table would be updated to match the `EMPLOYEES` table.

```
SELECT *
FROM COPY_EMP;
```

no rows selected

```
MERGE INTO copy_emp c
    USING employees e
    ON (c.employee_id = e.employee_id)
WHEN MATCHED THEN
    UPDATE SET
        ...
WHEN NOT MATCHED THEN
    INSERT VALUES...;
```

```
SELECT *
FROM COPY_EMP;
```

20 rows selected.

Summary

| Statement | Description |
|-----------|--|
| INSERT | Adds a new row to the table |
| UPDATE | Modifies existing rows in the table |
| DELETE | Removes existing rows from the table |
| MERGE | Conditionally inserts or updates data in a table |